

## Interactive / complex / 7

query	Interactive / complex / 7																																												
title	Recent likers																																												
pattern	<pre>graph LR     P1[Person id = \$id] -- knows (?) --&gt; P2[person: Person id firstName lastName]     P1 -- hasCreator --&gt; M[Message id content/imageFile]     M -- likes --&gt; P2</pre>																																												
desc.	Given a start Person, find (most recent) Likes on any of start Person’s Messages. Find Persons that Liked any of start Person’s Messages, the Messages they liked most recently, creation date of that Like, and the latency (in minutes) between creation of Messages and Like. Additionally, for each Person found return a flag indicating whether the liker is a friend of start Person. In the case that a Person Liked multiple Messages at the same time, return the Message with lowest identifier.																																												
params	<table><tr><td>1</td><td>Person.id</td><td>64-bit Integer</td><td></td></tr></table>					1	Person.id	64-bit Integer																																					
1	Person.id	64-bit Integer																																											
result	<table><tr><td>1</td><td>Person.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>2</td><td>Person.firstName</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>Person.lastName</td><td>String</td><td>R</td><td></td></tr><tr><td>4</td><td>Like.creationDate</td><td>DateTime</td><td>R</td><td></td></tr><tr><td>5</td><td>Message.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>6</td><td>Message.content or Post.imageFile</td><td>String</td><td>R</td><td></td></tr><tr><td>7</td><td>latency</td><td>32-bit Integer</td><td>C</td><td>Duration between creation of Message and Like, in minutes</td></tr><tr><td>8</td><td>isNew</td><td>Boolean</td><td>C</td><td>false if liker Person is friend of start Person, true otherwise</td></tr></table>					1	Person.id	ID	R		2	Person.firstName	String	R		3	Person.lastName	String	R		4	Like.creationDate	DateTime	R		5	Message.id	ID	R		6	Message.content or Post.imageFile	String	R		7	latency	32-bit Integer	C	Duration between creation of Message and Like, in minutes	8	isNew	Boolean	C	false if liker Person is friend of start Person, true otherwise
1	Person.id	ID	R																																										
2	Person.firstName	String	R																																										
3	Person.lastName	String	R																																										
4	Like.creationDate	DateTime	R																																										
5	Message.id	ID	R																																										
6	Message.content or Post.imageFile	String	R																																										
7	latency	32-bit Integer	C	Duration between creation of Message and Like, in minutes																																									
8	isNew	Boolean	C	false if liker Person is friend of start Person, true otherwise																																									
sort	<table><tr><td>1</td><td>Like.creationDate</td><td>↓</td><td></td></tr><tr><td>2</td><td>Person.id</td><td>↑</td><td></td></tr></table>					1	Like.creationDate	↓		2	Person.id	↑																																	
1	Like.creationDate	↓																																											
2	Person.id	↑																																											
limit	20																																												
CPs	2.2, 2.3, 3.3, 5.1																																												
relevance	This query looks for paths of length two, starting from a given Person, moving to its published messages and then to Persons who liked them. It tests several aspects related to join optimization, both at query optimization plan level and execution engine level. On the one hand, many of the columns needed for the projection are only needed in the last stages of the query, so the optimizer is expected to delay the projection until the end. This query implies accessing 2-hop data, and as a consequence, index accesses are expected to be scattered. We expect to observe variate cardinalities, depending on the characteristics of the input parameter, so properly selecting the join operators will be crucial. This query has a lot of correlated sub-queries, so it is testing the ability to flatten the query execution plans.																																												