

LDBC Social Network Benchmark (SNB) – 0.3.0

ABSTRACT

LDBC's Social Network Benchmark (LDBC SNB) is an effort intended to test various functionalities of systems used for graph-like data management. For this, LDBC SNB uses the recognizable scenario of operating a social network, characterized by its graph-shaped data.

LDBC SNB consists of two workloads that focus on different functionalities: the interactive workload (interactive transactional queries) and the business intelligence workload (analytical queries).

This document contains the definition of the Interactive Workload and the first draft of the Business Intelligence Workload. This includes a detailed explanation of the data used in the LDBC SNB benchmark, a detailed description for all queries, and instructions on how to generate the data and run the benchmark with the provided software.

EXECUTIVE SUMMARY

The new data economy era, based on complexly structured, distributed and large datasets, has brought on new demands on data management and analytics. As a consequence, new industry actors have appeared, offering technologies specially built for the management of graph-like data. Also, traditional database technologies, such as relational databases, are being adapted to the new demands to remain competitive.

LDBC's Social Network Benchmark (LDBC SNB) is an industrial and academic initiative, formed by principal actors in the field of graph-like data management. Its goal is to define a framework where different graph based technologies can be fairly tested and compared, that can drive the identification of systems' bottlenecks and required functionalities, and can help researchers to open new research frontiers.

The philosophy around which LDBC SNB is designed is to be easy to understand, flexible and cheap to adopt. For all these reasons, LDBC SNB will propose different workloads representing all the usage scenarios of graph-like database technologies, hence, targeting systems of different nature and characteristics. In order increase its adoption by industry and research institutions, LDBC SNB provides all necessary software, which are designed to be easy to use and deploy at a small cost.

This document contains:

- A detailed specification of the data used in the whole LDBC SNB benchmark.
- A detailed specification of the workloads.
- A detailed specification of the execution rules of the benchmark.
- A detailed specification of the auditing rules and the full disclosure report's required contents.

TABLE OF CONTENTS

1	INTRODUCTION	8
1.1	Motivation for the Benchmark	8
1.2	Relevance to the Industry	8
1.3	General Benchmark Overview	9
1.4	Related Projects	10
1.5	Participation of Industry and Academia	10
2	BENCHMARK SPECIFICATION	12
2.1	Requirements	12
2.2	Software and Useful Links	12
2.3	Data	12
2.3.1	Data Types	12
2.3.2	Data Schema	13
2.3.3	Output Data	17
2.3.4	Scale Factors	19
3	WORKLOADS	20
3.1	Query Description Format	20
3.2	Query Definitions	20
3.2.1	Notation	20
3.2.2	Morphisms	20
3.2.3	Graph Patterns	21
3.3	Substitution Parameters	21
3.4	Load Definition	22
4	INTERACTIVE WORKLOAD	24
4.1	Complex Reads Query Descriptions	24
4.2	Short Reads Query Descriptions	39
4.3	Update Query Descriptions	43
5	BUSINESS INTELLIGENCE WORKLOAD	48
5.1	Read Query Descriptions	48
5.2	Update Query Descriptions	72
6	AUDITING RULES	73
6.1	Preparation	73
6.1.1	Collect System Details	73
6.1.2	Setup the Benchmark Environment	73
6.1.3	Load Data	74
6.2	Running the Benchmark	74
6.3	Recovery	74
6.4	Serializability	75
7	RELATED WORK	76
7.1	Graph Database Benchmarks	76
7.2	Scalable Graph Generators	76
7.3	LDBC Publications	76

A	CHOKE POINTS	78
A.1	Aggregation Performance	78
A.2	Join Performance	80
A.3	Data Access Locality	81
A.4	Expression Calculation	82
A.5	Correlated Sub-queries	82
A.6	Parallelism and Concurrency	83
A.7	RDF and Graph Specifics	84
B	SCALE FACTOR STATISTICS	85
B.1	Scale Factor Statistics	85

ACKNOWLEDGMENTS

Special thanks to all the people that have contributed to the development of this benchmark:

- Renzo Angles (Universidad de Talca)
- Alex Averbuch (Neo Technologies)
- Peter Boncz (Vrije Universiteit Amsterdam and CWI)
- Orri Erling (Google)
- Andrey Gubichev (Google)
- Moritz Kaufmann (Tableau)
- Josep Lluís Larriba Pey (Universitat Politècnica de Catalunya)
- Minh-Duc Pham (Altran)
- Marcus Paradies (SAP)
- Arnau Prat (Sparsity Technologies)
- Mirko Spasić (Openlink Software)
- Norbert Martínez (Huawei Technologies)
- Gábor Szárnyas (MTA-BME Lendület Research Group on Cyber-Physical Systems)

DEFINITIONS

DATAGEN: Is the data generator provided by the LDBC SNB, which is responsible for generating the data needed to run the benchmark.

DBMS: A DataBase Management System.

LDBC SNB: Linked Data Benchmark Council Social Network Benchmark.

Query Mix: Refers to the ratio between read and update queries of a workload, and the frequency at which they are issued.

SF (Scale Factor): The LDBC SNB is designed to target systems of different size and scale. The scale factor determines the size of the data used to run the benchmark, measured in Gigabytes.

SUT: The System Under Test is defined to be the database system where the benchmark is executed.

Test Driver: A program provided by the LDBC SNB, which is responsible for executing the different workloads and gathering the results.

Full Disclosure Report (FDR): The FDR is a document which allows reproduction of any benchmark result by a third party. This contains complete description of the SUT and the circumstances of the benchmark run, e.g. configuration of SUT, dataset and test driver, etc.

Test Sponsor: The Test Sponsor is the company officially submitting the Result with the FDR and will be charged the filing fee. Although multiple companies may sponsor a Result together, for the purposes of the LDBC processes the Test Sponsor must be a single company. A Test Sponsor need not be a LDBC member. The Test Sponsor is responsible for maintaining the FDR with any necessary updates or corrections. The Test Sponsor is also the name used to identify the Result.

Workload: A workload refers to a set of queries of a given nature (i.e. interactive, analytical, business), how they are issued and at which rate.

1 INTRODUCTION

1.1 Motivation for the Benchmark

The new era of data economy, based on large, distributed, and complexly structured data sets, has brought on new and complex challenges in the field of data management and analytics. These data sets, usually modeled as large graphs, have attracted both industry and academia, due to new opportunities in research and innovation they offer. This situation has also opened the door for new companies to emerge, offering new non-relational and graph-like technologies that are called to play a significant role in upcoming years.

The change in the data paradigm calls for new benchmarks to test these new emerging technologies, as they set a framework where different systems can compete and be compared in a fair way, they let technology providers identify the bottlenecks and gaps of their systems and, in general, drive the research and development of new information technology solutions. Without them, the uptake of these technologies is at risk by not providing the industry with clear, user-driven targets for performance and functionality.

The LDBC Social Network Benchmark (LDBC SNB) aims at being a comprehensive benchmark by setting the rules for the evaluation of graph-like data management technologies. LDBC SNB is designed to be a plausible look-alike of all the aspects of operating a social network site, as one of the most representative and relevant use cases of modern graph-like applications.

LDBC SNB includes the Interactive Workload [4], which consists of user-centric transactional-like interactive queries, and the Business Intelligence Workload, which includes analytic queries to respond to business-critical questions. Initially, a graph analytics workload was also included in the roadmap of LDBC SNB, but this was finally delegated to the Graphalytics benchmark project [6], which was adopted as an official LDBC graph analytics benchmark. LDBC SNB and Graphalytics combined target a broad range of systems with different nature and characteristics. LDBC SNB and Graphalytics aim at capturing the essential features of these scenarios while abstracting away details of specific business deployments.

This document contains the definition of the Interactive Workload and the first draft of the Business Intelligence Workload. This includes a detailed explanation of the data used in the LDBC SNB benchmark, a detailed description for all queries, and instructions on how to generate the data and run the benchmark with the provided software.

1.2 Relevance to the Industry

LDBC SNB is intended to provide the following value to different stakeholders:

- For **end users** facing graph processing tasks, LDBC SNB provides a recognizable scenario against which it is possible to compare merits of different products and technologies. By covering a wide variety of scales and price points, LDBC SNB can serve as an aid to technology selection.
- For **vendors** of graph database technology, LDBC SNB provides a checklist of features and performance characteristics that helps in product positioning and can serve to guide new development.
- For **researchers**, both industrial and academic, the LDBC SNB dataset and workload provide interesting challenges in multiple choke-point areas, such as query optimization, (distributed) graph analysis, transactional throughput, and provides a way to objectively compare the effectiveness and efficiency of new and existing technology in these areas.

The technological scope of LDBC SNB comprises all systems that one might conceivably use to perform social network data management tasks:

- **Graph database systems** (e.g. Neo4j, InfiniteGraph, Sparksee, Titan) are novel technologies aimed at storing directed and labeled graphs. They support graph traversals, typically by means of APIs, though some of them also support dedicated graph-oriented query languages (e.g. Neo4j's Cypher). These systems' internal structures are typically designed to store dynamic graphs that change over time. They offer

support for transactional queries with some degree of consistency, and value-based indexes to quickly locate nodes and edges. Finally, their architecture is typically single-machine (non-cluster). These systems can potentially implement all three workloads, though Interactive and Business Intelligence workloads are where they will presumably be more competitive.

- **Graph processing frameworks** (e.g. Giraph, Signal/Collect, GraphLab, Green Marl) are designed to perform global graph computations, executed in parallel or in a lockstep fashion. These computations are typically long latency, involving many nodes and edges and often consist of approximation answers to NP-complete problems. These systems expose an API, sometimes following a vertex-centric paradigm, and their architecture targets both single-machine and cluster systems. These systems will likely implement the Graph Analytics workload.
- **RDF database systems** (e.g. OWLIM, Virtuoso, BigData, Jena TDB, Stardog, Allegrograph) are systems that implement the SPARQL 1.1 query language, similar in complexity to SQL-92, which allows for structured queries, and simple traversals. RDF database systems often come with additional support for simple reasoning (sameAs, subClass), text search, and geospatial predicates. RDF database systems generally support transactions, but not always with full concurrency and serializability and their supposed strength is integrating multiple data sources (e.g. DBpedia). Their architecture is both single-machine and clustered, and they will likely target Interactive and Business Intelligence workloads.
- **Relational database systems** (e.g. Postgres, MySQL, Oracle, IBM DB2, Microsoft SQL Server, Virtuoso, MonetDB, Vectorwise, Vertica, but also Hive and Impala) treat graph data relationally, and queries are formulated in SQL and/or PL/SQL. Both single-machine and cluster systems exist. They do not normally support recursion or stateful recursive algorithms, which makes them not at home in the Graph Analytics workloads.
- **NoSQL database systems** (e.g. key-value stores such as HBase, REDIS, MongoDB, CouchDB, or even MapReduce systems like Hadoop and Pig) are cluster-based and scalable. Key-value stores could possibly implement the Interactive Workload, though its navigational aspects would pose some problems as potentially many key-value lookups are needed. MapReduce systems could be suited for the Graph Analytics workload, but their query latency would presumably be so high that the Business Intelligence workload would not make sense, though we note that some of the key-value stores (e.g. MongoDB) provide a MapReduce query functionality on the data that it stores which could make it suited for the Business Intelligence workload.

1.3 General Benchmark Overview

LDBC SNB aims at being a complete benchmark, designed with the following goals in mind:

- **Rich coverage.** LDBC SNB is intended to cover most demands encountered in the management of complexly structured data.
- **Modularity.** LDBC SNB is broken into parts that can be individually addressed. In this manner LDBC SNB stimulates innovation without imposing an overly high threshold for participation.
- **Reasonable implementation cost.** For a product offering relevant functionality, the effort for obtaining initial results with SNB should be small, in the order of days.
- **Relevant selection of challenges.** Benchmarks are known to direct product development in certain directions. LDBC SNB is informed by the state-of-the-art in database research so as to offer optimization challenges for years to come while not having a prohibitively high threshold for entry.
- **Reproducibility and documentation of results.** LDBC SNB will specify the rules for full disclosure of benchmark execution and for auditing of benchmark runs. The workloads may be run on any equipment but the exact configuration and price of the hardware and software must be disclosed.

LDBC SNB benchmark is modeled around the operation of a real social network site. A social network site represents a relevant use case for the following reasons:

- It is simple to understand for a large audience, as it is arguably present in our every-day life in different shapes and forms.
- It allows testing a complete range of interesting challenges, by means of different workloads targeting systems of different nature and characteristics.
- A social network can be scaled, allowing the design of a scalable benchmark targeting systems of different sizes and budgets.

In Section 2.3, LDBC SNB defines the schema of the data used in the benchmark. The schema represents a realistic social network, including people and their activities in the social network during a period of time. Personal information of each person, such as name, birthday, interests or places where people work or study, is included. A persons' activity is represented in the form of friendship relationships and content sharing (i.e. messages and pictures). LDBC SNB provides a scalable synthetic data generator based on the MapReduce paradigm, which produces networks with the described schema with distributions and correlations similar to those expected in a real social network. Furthermore, the data generator is designed to be user-friendly. The proposed data schema is shared by all the different proposed workloads, those we currently have, and those that will be proposed in the future.

In chapter 3, the Interactive Workload and the first draft of the Business Intelligence workload are proposed. Workloads are designed to mimic the different usage scenarios found in operating a real social network site, and each of them targets one or more types of systems. Each workload defines a set of queries and query mixes, designed to stress the SUTs in different choke-point areas, while being credible and realistic. The Interactive workload reproduces the interaction between the users of the social network by including lookups and transactions, which update small portions of the database. These queries are designed to be interactive and target systems capable of responding to such queries with low latency for multiple concurrent users. The Business Intelligence workload represents analytic queries a social network company would like to perform in the social network, to take advantage of the data and to discover new business opportunities. This workload explores moderate to large portions of the graph from different entities, and performs more resource-intensive operations.

LDBC SNB provides an execution test driver, which is responsible for executing the workloads and gathering the results. The driver is designed with simplicity and portability in mind to ease the implementation on systems with different nature and characteristics at a low implementation cost. Furthermore, it automatically handles the validation of the queries by means of a validation dataset provided by LDBC. The overall philosophy of LDBC SNB is to provide the necessary software tools to run the benchmark, and therefore to reduce the benchmark's entry point as much as possible.

1.4 Related Projects

Along the Social Network Benchmark, LDBC [1] provides other benchmarks as well:

- The Semantic Publishing Benchmark (SPB)¹ measures the performance of RDF engines operating on semantic data sets.
- The Graphalytics benchmark [7] measures the performance of graph analysis operations (e.g. PageRank, local clustering coefficient).

1.5 Participation of Industry and Academia

The list of institutions that take part in the definition and development of LDBC SNB is formed by relevant actors from both the industry and academia in the field of linked data management. All the participants have contributed with their experience and expertise in the field, making a credible and relevant benchmark, which meets all the desired needs. The list of participants is the following:

- FOUNDATION FOR RESEARCH AND TECHNOLOGY HELLAS

¹<http://ldbpcouncil.org/benchmarks/spb>

- MTA-BME LENDUELET RESEARCH GROUP ON CYBER-PHYSICAL SYSTEMS
- NEO4J
- ONTOTEXT
- OPENLINK
- TECHNISCHE UNIVERSITAET MUENCHEN
- UNIVERSITAET INNSBRUCK
- UNIVERSITAT POLITECNICA DE CATALUNYA
- VRIJE UNIVERSITEIT AMSTERDAM

Besides the aforementioned institutions, during the development of the benchmark several meetings with the technical and users community have been conducted, receiving an invaluable feedback that has contributed to the whole development of the benchmark in every of its aspects.

2 BENCHMARK SPECIFICATION

2.1 Requirements

LDBC SNB is designed to be flexible and to have an affordable entry point. From small single node and in memory systems to large distributed multi-node clusters have its own place in LDBC SNB. Therefore, the requirements to fulfill for executing LDBC SNB are limited to pure software requirements to be able to run the tools. All the software provided by LDBC SNB have been developed and tested under Linux-based operating systems.

LDBC SNB does not impose the usage of any specific type of system, as it targets systems of different nature and characteristics, from graph databases, graph processing frameworks and RDF systems, to traditional relational database management systems. Consequently, any language or API capable of expressing the proposed queries can be used. Similarly, data can be stored in the most convenient manner the test sponsor may decide.

2.2 Software and Useful Links

- **LDBC Driver 0.3** – https://github.com/ldbc/ldbc_driver: The driver responsible for executing the LDBC SNB workload.
- **DATAGEN 0.2.5** – https://github.com/ldbc/ldbc_snb_datagen: The data generator used to generate the datasets of the benchmark.

2.3 Data

This section introduces the data used by LDBC SNB. This includes the different data types, the data schema, how it is generated and the different scale factors.

2.3.1 Data Types

Table 2.1 describes the different types used in the whole benchmark.

Type	Description
ID	integer type with 64-bit precision. All IDs within a single entity type (e.g. Post) are unique, but different entity types (e.g. a Forum and a Post) might have the same ID.
32-bit Integer	integer type with 32-bit precision
64-bit Integer	integer type with 64-bit precision
String	variable length text of size 40 Unicode characters
Long String	variable length text of size 256 Unicode characters
Text	variable length text of size 2000 Unicode characters
Date	date with a precision of a day, encoded as a string with the following format: <i>yyyy-mm-dd</i> , where <i>yyyy</i> is a four-digit integer representing the year, the year, <i>mm</i> is a two-digit integer representing the month and <i>dd</i> is a two-digit integer representing the day.
DateTime	date with a precision of milliseconds, encoded as a string with the following format: <i>yyyy-mm-ddTHH:MM:ss.sss+0000</i> , where <i>yyyy</i> is a four-digit integer representing the year, the year, <i>mm</i> is a two-digit integer representing the month and <i>dd</i> is a two-digit integer representing the day, <i>HH</i> is a two-digit integer representing the hour, <i>MM</i> is a two digit integer representing the minute and <i>ss.sss</i> is a five digit fixed point real number representing the seconds up to millisecond precision. Finally, the <i>+0000</i> of the end represents the timezone, which in this case is always GMT.

Table 2.1: Description of the data types.

2.3.2 Data Schema

Figure 2.1 shows the data schema in UML. The schema defines the structure of the data used in the benchmark in terms of entities and their relations. Data represents a snapshot of the activity of a social network during a period of time. Data includes entities such as Persons, Organisations, and Places. The schema also models the way persons interact, by means of the friendship relations established with other persons, and the sharing of content such as messages (both textual and images), replies to messages and likes to messages. People form groups to talk about specific topics, which are represented as tags.

LDBC SNB has been designed to be flexible and to target systems of different nature and characteristics. As such, it does not force any particular internal representation of the schema. The DATAGEN component supports multiple output data formats to fit the needs of different types of systems, including RDF, relational DBMS and graph DBMS.

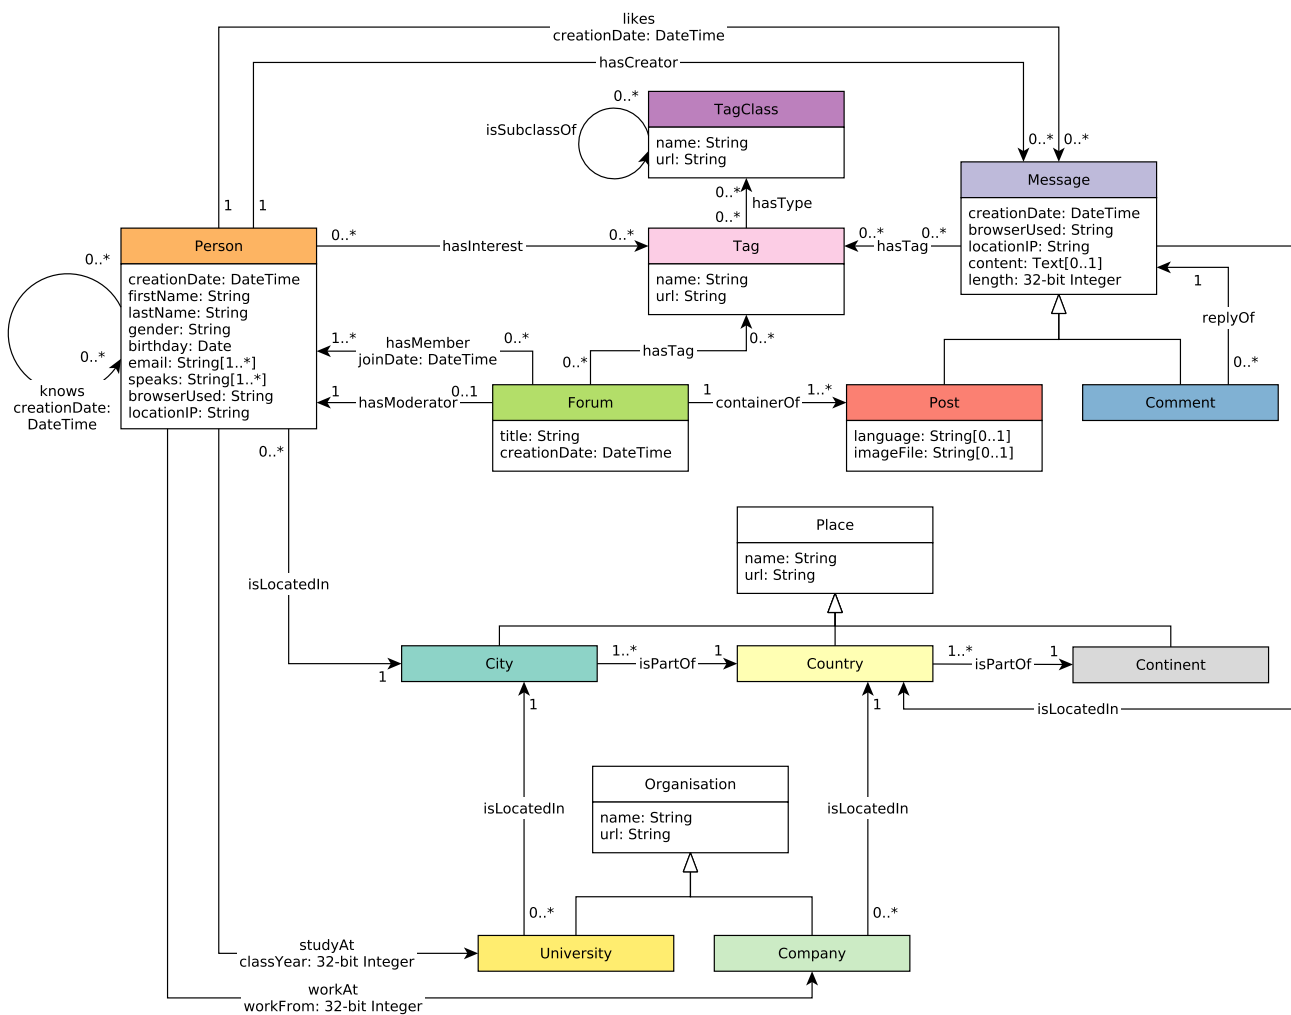


Figure 2.1: The LDBC SNB data schema

The schema specifies different entities, their attributes and their relations. All of them are described in the following sections.

Textual Restrictions and Notes

- Posts have content or imageFile. They have one of them but not both. The one they do not have is an empty string.
- Posts in a forum can be created by a non-member person if and only if that person is a moderator.

Entities

City: a sub-class of a Place, and represents a city of the real world. City entities are used to specify where persons live, as well as where universities operate.

Comment: a sub-class of a Message, and represents a comment made by a person to an existing message (either a Post or a Comment).

Company: a sub-class of an Organisation, and represents a company where persons work.

Country: a sub-class of a Place, and represents a continent of the real world.

Forum: a meeting point where people post messages. Forums are characterized by the topics (represented as tags) people in the forum are talking about. Although from the schema's perspective it is not evident, there exist three different types of forums: persons' personal walls, image albums, and groups. They are distinguished by their titles. Table 2.2 shows the attributes of Forum entity.

Attribute	Type	Description
id	ID	The identifier of the forum.
title	Long String	The title of the forum.
creationDate	DateTime	The date the forum was created.

Table 2.2: Attributes of Forum entity.

Message: an abstract entity that represents a message created by a person. Table 2.3 shows the attributes of Message abstract entity.

Attribute	Type	Description
id	ID	The identifier of the message.
browserUsed	String	The browser used by the Person to create the message.
creationDate	DateTime	The date the message was created.
locationIP	String	The IP of the location from which the message was created.
content	Text[0..1]	The content of the message.
length	32-bit Integer	The length of the content.

Table 2.3: Attributes of Message interface.

Organisation: an institution of the real world. Table 2.4 shows the attributes of Organisation entity.

Attribute	Type	Description
id	ID	The identifier of the organisation.
name	Long String	The name of the organisation.
url	Long String	The URL of the organisation.
url	Long String	The URL of the organization.

Table 2.4: Attributes of Organisation entity.

Person: the avatar a real world person creates when he/she joins the network, and contains various information about the person as well as network related information. Table 2.5 shows the attributes of Person entity.

Attribute	Type	Description
id	ID	The identifier of the person.
firstName	String	The first name of the person.
lastName	String	The last name of the person.
gender	String	The gender of the person.
birthDay	Date	The birthday of the person.
email	Long String[1..*]	The set of emails the person has.
speaks	String[1..*]	The set of languages the person speaks.
browserUser	String	The browser used by the person when he/she registered to the social network.
locationIp	String	The IP of the location from which the person was registered to the social network.
creationDate	DateTime	The date the person joined the social network.

Table 2.5: Attributes of Person entity.

Place: a place in the world. Table 2.6 shows the attributes of Place entity.

Attribute	Type	Description
id	ID	The identifier of the place.
name	Long String	The name of the place.
url	Long String	The URL of the place.

Table 2.6: Attributes of Place entity.

Post: a sub-class of Message, that is posted in a forum. Posts are created by persons into the forums where they belong. Posts contain either content or imageFile, always one of them but never both. The one they do not have is an empty string. Table 2.7 shows the attributes of Post entity.

Attribute	Type	Description
language	String[0..1]	The language of the post.
imageFile	String[0..1]	The image file of the post.

Table 2.7: Attributes of Post entity.

Tag: a topic or a concept. Tags are used to specify the topics of forums and posts, as well as the topics a person is interested in. Table 2.8 shows the attributes of Tag entity.

Attribute	Type	Description
id	ID	The identifier of the tag.
name	Long String	The name of the tag.
url	Long String	The URL of the tag.

Table 2.8: Attributes of Tag entity.

TagClass: a class or a category used to build a hierarchy of tags. Table 2.9 shows the attributes of TagClass entity.

Attribute	Type	Description
id	ID	The identifier of the tagclass.
name	Long String	The name of the tagclass.
url	Long String	The URL of the tagclass.

Table 2.9: Attributes of TagClass entity.

University: a sub-class of Organisation, and represents an institution where persons study.

Relations

Relations connect entities of different types. Entities are defined by their "id" attribute.

Name	Tail	Head	Type	Description						
containerOf	Forum[1]	Post[1..*]	D	A Forum and a Post contained in it						
hasCreator	Message[0..*]	Person[1]	D	A Message and its creator (Person)						
hasInterest	Person[0..*]	Tag[0..*]	D	A Person and a Tag representing a topic the person is interested in						
hasMember	Forum[0..*]	Person[1..*]	D	A Forum and a member (Person) of the forum <table><tr><td>Attribute</td><td>joinDate</td></tr><tr><td>Type</td><td>DateTime</td></tr><tr><td>Description</td><td>The Date the person joined the forum</td></tr></table>	Attribute	joinDate	Type	DateTime	Description	The Date the person joined the forum
Attribute	joinDate									
Type	DateTime									
Description	The Date the person joined the forum									
hasModerator	Forum[0..*]	Person[1]	D	A Forum and its moderator (Person)						
hasTag	Message[0..*]	Tag[0..*]	D	A Message and a Tag representing the message's topic						
hasTag	Forum[0..*]	Tag[0..*]	D	A Forum and a Tag representing the forum's topic						
hasType	Tag[0..*]	TagClass[0..*]	D	A Tag and a TagClass the tag belongs to						
isLocatedIn	Company[0..*]	Country[1]	D	A Company and its home Country						
isLocatedIn	Message[0..*]	Country[1]	D	A Message and the Country from which it was issued						
isLocatedIn	Person[0..*]	City[1]	D	A Person and their home City						
isLocatedIn	University[0..*]	City[1]	D	A University and the City where the university is						
isPartOf	City[1..*]	Country[1]	D	A City and the Country it is part of						
isPartOf	Country[1..*]	Continent[1]	D	A Country and the Continent it is part of						
isSubclassOf	TagClass[0..*]	TagClass[0..*]	D	A TagClass its parent TagClass						
knows	Person[0..*]	Person[0..*]	U	Two Persons that know each other <table><tr><td>Attribute</td><td>creationDate</td></tr><tr><td>Type</td><td>DateTime</td></tr><tr><td>Description</td><td>The date the knows relation was established</td></tr></table>	Attribute	creationDate	Type	DateTime	Description	The date the knows relation was established
Attribute	creationDate									
Type	DateTime									
Description	The date the knows relation was established									
likes	Person[0..*]	Message[0..*]	D	A Person that likes a Message <table><tr><td>Attribute</td><td>creationDate</td></tr><tr><td>Type</td><td>DateTime</td></tr><tr><td>Description</td><td>The date the like was issued</td></tr></table>	Attribute	creationDate	Type	DateTime	Description	The date the like was issued
Attribute	creationDate									
Type	DateTime									
Description	The date the like was issued									

replyOf	Comment[0..*]	Message[1]	D	A Comment and the Message it replies
studyAt	Person[0..*]	University[0..*]	D	A Person and a University it has studied
				Attribute classYear
				Type 32-bit Integer
				Description The year the person graduated
workAt	Person[0..*]	Company[0..*]	D	A Person and a Company it works
				Attribute workFrom
				Type 32-bit Integer
				Description The year the person started to work at that company

Table 2.10: Description of the data relations.

Domain Concepts

A *thread* consist of Messages, starting with a single Post and Comments that transitively reply to that Post.

2.3.3 Output Data

DATAGEN produces outputs three different items:

- **Dataset:** The dataset to be bulk loaded by the SUT. It corresponds to roughly the 90% of the total generated network.
- **Update Streams:** A set of update streams containing update queries, which are used by the driver to generate the update queries of the workloads. This update streams correspond to the remaining 10% of the generated dataset.
- **Substitution Parameters:** A set of files containing the different parameter bindings that will be used by the driver to generate the read queries of the workloads.

The SUT have to take care only of the generated Dataset to be bulk loaded. Three different formats are supported by DATAGEN:

- **CSV:** Data output in CSV format, one file per different entity and on file per different relation. Also, there is a file por those attributes whose cardinality is larger than one (i.e. Person.email, Person.speaks, etc.).
- **CSVMergeForeign:** Similar to CSV format, but in this case, those relations of the form 1 to 1 and 1 to N, are stored in the tail entity file as a foreign keys.
- **Turtle:** Dataset in Turtle format for RDF systems.

CSV

This is a comma separated format. Each entity, relation and properties with a cardinality larger than one, are output in a separate file. Generated files are summarized at Table 2.11. Depending on the number of threads used for generating the dataset, the number of files varies, since there is a file generated per thread. The * in the file names indicates a number between 0 and NumberOfThreads – 1.

CSV_MERGE_FOREIGN

This is a comma separated format. It is similar to CSV, but those relations connecting two entities A and B, where an entity A has a cardinality of one, A is output as a column of entity B. Generated files are summarized at Table 2.12. Depending on the number of threads used for generating the dataset, the number of files

File	Content
comment_*.csv	id creationDate locationIP browserUsed content length
comment_hasCreator_person_*.csv	Comment.id Person.id
comment_hasTag_tag_*.csv	Comment.id Tag.id
comment_isLocatedIn_place_*.csv	Comment.id Place.id
comment_replyOf_comment_*.csv	Comment.id Comment.id
comment_replyOf_post_*.csv	Comment.id Post.id
forum_*.csv	id title creationDate
forum_containerOf_post_*.csv	Forum.id Post.id
forum_hasMember_person_*.csv	Forum.id Person.id joinDate
forum_hasModerator_person_*.csv	Forum.id Person.id
forum_hasTag_tag_*.csv	Forum.id Tag.id
organisation_*.csv	id type("university", "company") name url
organisation_isLocatedIn_place_*.csv	Organisation.id Place.id
person_*.csv	id firstName lastName gender birthday creationDate locationIP browserUsed
person_email_emailaddress_*.csv	Person.id email
person_hasInterest_tag_*.csv	Person.id Tag.id
person_isLocatedIn_place_*.csv	Person.id Place.id
person_knows_person_*.csv	Person.id Person.id creationDate
person_likes_comment_*.csv	Person.id Post.id creationDate
person_likes_post_*.csv	Person.id Post.id creationDate
personSpeaks_language_*.csv	Person.id language
person_studyAt_organisation_*.csv	Person.id Organisation.id classYear
person_workAt_organisation_*.csv	Person.id Organisation.id workFrom
place_*.csv	id name url type("city", "country", "continent")
place_isPartOf_place_*.csv	Place.id Place.id
post_*.csv	id imageFile creationDate locationIP browserUsed language content length
post_hasCreator_person_*.csv	Post.id Person.id
post_hasTag_tag_*.csv	Post.id Tag.id
post_isLocatedIn_place.csv	Post.id Place.id
tag_*.csv	id name url
tag_hasType_tagclass_*.csv	Tag.id TagClass.id
tagclass_*.csv	id name url
tagclass_isSubclassOf_tagclass_*.csv	TagClass.id TagClass.id

Table 2.11: Files output by the CSV serializer (33 in total)

varies, since there is a file generated per thread. The * in the file names indicates a number between 0 and NumberOfThreads – 1.

File	Content
comment_*.csv	id creationDate locationIP browserUsed content length creator place replyOfPost replyOfComment
comment_hasTag_tag_*.csv	Comment.id Tag.id
forum_*.csv	id title creationDate moderator
forum_hasMember_person_*.csv	Forum.id Person.id joinDate
forum_hasTag_tag_*.csv	Forum.id Tag.id
organisation_*.csv	id type("university", "company") name url
person_*.csv	id firstName lastName gender birthday creationDate locationIP browserUsed place
person_email_emailaddress_*.csv	Person.id email
person_hasInterest_tag_*.csv	Person.id Tag.id
person_knows_person_*.csv	Person.id Person.id creationDate
person_likes_comment_*.csv	Person.id Post.id creationDate
person_likes_post_*.csv	Person.id Post.id creationDate
personSpeaks_language_*.csv	Person.id language
person_studyAt_organisation_*.csv	Person.id Organisation.id classYear
person_workAt_organisation_*.csv	Person.id Organisation.id workFrom
place_*.csv	id name url type("city", "country", "continent")
post_*.csv	id imageFile creationDate locationIP browserUsed language content length creator Forum.id place
post_hasTag_tag_*.csv	Post.id Tag.id
tag_*.csv	id name url
tag_hasType_tagclass_*.csv	Tag.id TagClass.id
tagclass_*.csv	id name url
tagclass_isSubclassOf_tagclass_*.csv	TagClass.id TagClass.id

Table 2.12: Files output by the CSV_MERGE_FOREIGN serializer (22 in total)

Turtle

This is the standard Turtle¹ format. DATAGEN outputs two files: `0_ldbc_socialnet_static_dbp.ttl` and `0_ldbc_socialnet.ttl`.

2.3.4 Scale Factors

LDBC SNB defines a set of scale factors (SFs), targeting systems of different sizes and budgets. SFs are computed based on the ASCII size in Gigabytes of the generated output files using the CSV serializer. For example, SF 1 weights roughly 1 GB in CSV format, SF 3 weights roughly 3 GB and so on and so forth. The proposed SFs are the following: 1, 3, 10, 30, 100, 300, 1000. The Test Sponsor may select the SF that better fits their needs, by properly configuring the DATAGEN.

The size of the resulting dataset, is mainly affected by the following configuration parameters: the number of persons and the number of years simulated. Different SFs are computed by scaling the number of Persons in the network, while fixing the number of years simulated. Table 2.13 shows the parameters used in each of the SFs.

Scale Factor	1	3	10	30	100	300	1000
# of Persons	11K	27K	73K	182K	499K	1.25M	3.6M
# of Years	3	3	3	3	3	3	3
Start Year	2010	2010	2010	2010	2010	2010	2010

Table 2.13: Parameters of each scale factor.

For example, SF 30 consists of the activity of a social network of 182K users during a period of three years, starting from 2010.

¹Description of the Turtle RDF format <http://www.w3.org/TR/turtle/>

3 WORKLOADS

3.1 Query Description Format

Queries are described in natural language using a well-defined structure that consists of three sections: *description*, a concise textual description of the query; *parameters*, a list of input parameters and their types; and *results*, a list of expected results and their types. The syntax used in *parameters* and *results* sections is as follows:

- **Entity:** entity type in the dataset.
One word, possibly constructed by appending multiple words together, starting with uppercase character and following the camel case notation, e.g. TagClass represents an entity of type “TagClass”.
- **Relationship:** relationship type in the dataset.
One word, possibly constructed by appending multiple words together, starting with lowercase character and following the camel case notation, and surrounded by arrow to communicate direction, e.g. -worksAt-> represents a directed relationship of type “worksAt”.
- **Attribute:** attribute of an entity or relationship in the dataset.
One word, possibly constructed by appending multiple words together, starting with lowercase character and following the camel case notation, and prefixed by a “.” to dereference the entity/relationship, e.g. Person.firstName refers to “firstName” attribute on the “Person” entity, and -studyAt->.classYear refers to “classYear” attribute on the “studyAt” relationship.
- **Unordered Set:** an unordered collection of distinct elements.
Surrounded by { and } braces, with the element type between them, e.g. {String} refers to a set of strings.
- **Ordered List:** an ordered collection where duplicate elements are allowed.
Surrounded by [and] braces, with the element type between them, e.g. [String] refers to a list of strings.
- **Ordered Tuple:** a fixed length, fixed order list of elements, where elements at each position of the tuple have predefined, possibly different, types.
Surrounded by < and > braces, with the element types between them in a specific order e.g. <String, Boolean> refers to a 2-tuple containing a string value in the first element and a boolean value in the second, and [<String, Boolean>] is an ordered list of those 2-tuples.

Categorization of results. Results are categorized according to their source of origin:

- **Raw (R)**, if the result is returned with an unmodified value and type.
- **Calculated (C)**, if the result is calculated from other values and conditions.
- **Aggregated (A)**, if the result is an aggregated value, e.g. a count or a sum of another value. If a result is both calculated and aggregated (e.g. $\text{count}(x) + \text{count}(y)$ or $\text{avg}(x + y)$), it is considered an aggregated result.
- **Meta (M)**, if the result is based on type information, e.g. the type of the vertex.

3.2 Query Definitions

3.2.1 Notation

Intervals. Closed interval boundaries are denoted with [and], while open interval boundaries are denoted with (and). For example, $[\emptyset, 1)$ denoted an interval between 0 and 1, closed on the left and open on the right.

3.2.2 Morphisms

Discuss *homomorphism* vs. *isomorphism*

3.2.3 Graph Patterns

To illustrate queries, we use graph patterns such as Figure 3.1 with the following notation:




Figure 3.1: Example graph pattern.

- Vertices are marked as **entityName: EntityType** (camel case notation for both, starting with a lowercase character for the former and an uppercase character for the latter). The **entityName** is assigned if the vertex or its attributes are used later for aggregation or as part of the query result. If they are not used later, the **entityName:** part can be omitted.
- Positive conditions for edges are denoted with solid lines.
- Negative conditions for edges, i.e. edges that are not allowed in the graph, are denoted with **dashed red** lines.
- Edges without direction imply that there must be an edge in *at least one of the directions*.
- Filtering conditions are typeset in *italic*, e.g. *id = \$tag*.
- Attributes that should be returned are denoted in sans-serif font, e.g. name.
- Variable length paths, i.e. edges that can be traversed multiple times are denoted with *min...max, e.g. replyOf* or knows *1...2. By default, the value of min is 1, and the value of max is unlimited.
- Aggregations are shown in dashed boxes with the type of aggregation (count, sum, avg, etc.) in the upper right corner.

Keywords. The notation uses a small set of keywords:

- UNWIND unnests a list, i.e. produces a set of one-tuples. For example, UNWIND [1, 2, 3] results in {⟨1⟩, ⟨2⟩, ⟨3⟩}
- Aggregation operations: count, avg.
- Functions:
 - floor(x) (returns $\lfloor x \rfloor$),
 - year(date) (extracts the year from a given date),
 - month(date) (extracts the month from a given date).

Resolving ambiguity. Note that if the textual description and the graph pattern are different for a particular query (either due to an error or the lack of sophistication in the graphical syntax), *the textual description takes precedence*.

3.3 Substitution Parameters

Together with the dataset, DATAGEN produces a set of parameters per query type. Parameter generation is designed in such a way that for each query type, all of the generated parameters yield similar runtime behaviour of that query.

Specifically, the selection of parameters for a query template guarantees the following properties of the resulting queries:

- P1: the query runtime has a bounded variance: the average runtime corresponds to the behavior of the majority of the queries

- P2: the runtime distribution is stable: different samples of (e.g. 10) parameter bindings used in different query streams result in an identical runtime distribution across streams
- P3: the optimal logical plan (optimal operator order) of the queries is the same: this ensures that a specific query template tests the system's behavior under the well-chosen technical difficulty (e.g. handling voluminous joins or proper cardinality estimation for subqueries, etc.)

As a result, the amount of data that the query touches is roughly the same for every parameter binding, assuming that the query optimizer figures out a reasonable execution plan for the query. This is done to avoid bindings that cause unexpectedly long or short runtimes of queries, or even result in a completely different optimal execution plan. Such effects could arise due to the data skew and correlations between values in the generated dataset.

In order to get the parameter bindings for each of the queries, we have designed a *Parameter Curation* procedure that works in two stages:

1. for each query template for all possible parameter bindings, we determine the size of intermediate results in the *intended* query plan. Intermediate result size heavily influences the runtime of a query, so two queries with the same operator tree and similar intermediate result sizes at every level of this operator tree are expected to have similar runtimes. This analysis is effectively a side effect of data generation, that is we keep all the necessary counts (number of friends per user, number of posts of friends etc.) as we create the dataset.
2. then, a greedy algorithm selects (“curates”) those parameters with similar intermediate result counts from the domain of all the parameters.

Parameter bindings are stored in the `substitution_parameters` folder inside the data generator directory. Each query gets its bindings in a separate file. Every line of a parameter file is a JSON-formatted collection of key-value pairs (name of the parameter and its value). For example, the Query 1 parameter bindings are stored in file `query_1_param.txt`, and one of its lines may look like this:

```
{"PersonID": 1, "Name": "Lei", "PersonURI": "http://www.ldbc.eu/ldbc_socialnet/1.0/data/pers1"}
```

Depending on implementation, the SUT may refer to persons either by IDs (relational and graph databases) or URIs (RDF systems), so we provide both values for the Person parameter. Finally, parameters for short reads are taken from those in complex reads and updates.

3.4 Load Definition

LDBC SNB Test Driver is in charge of the execution of the Interactive Workload. At the beginning of the execution, the Test Driver creates a query mix by assigning to each query instance, a query issue time and a set of parameters taken from the generated substitution parameter set described above.

Query issue times have to be carefully assigned. Although substitution parameters are chosen in such a way that queries of the same type take similar time, not all query types have the same complexity and touch the same amount of data, which causes them to scale differently for the different scale factors. Therefore, if all query instances, regardless of their type, are issued at the same rate, those more complex queries will dominate the execution's result, making faster query types purposeless. To avoid this situation, each query type is executed at a different rate. The way the execution rate is decided, also depends on the nature of the query: complex read, short read or update.

Update queries' issue times are taken from the update streams generated by the data generator. These are the times where the actual event happened during the simulation of the social network. Complex reads' times are expressed in terms of update operations. For each complex read query type, a frequency value is assigned which specifies the relation between the number of updates performed per complex read. Table 3.1 shows the frequencies assigned to each query type for SF1. The frequencies of the different scale factors can be found in Section B.1.

Query Type	freq	Query Type	freq
Query 1	26	Query 8	45
Query 2	37	Query 9	157
Query 3	69	Query 10	30
Query 4	36	Query 11	16
Query 5	57	Query 12	44
Query 6	129	Query 13	19
Query 7	87	Query 14	49

Table 3.1: Frequencies for each query type for SF1.

Finally, short reads are inserted in order to balance the ratio between reads and writes, and to simulate the behavior of a real user of the social network. For each complex read instance, a sequence of short reads is planned. There are two types of short read sequences: Person centric and Message centric. Depending on the type of the complex read, one of them is chosen. Each sequence consists of a set of short reads which are issued in a row. The issue time assigned to each short read in the sequence is determined at run time, and is based on the completion time of the complex read it depends on. The substitution parameters for short reads are taken from the results of previously executed complex reads and short reads. Once a short read sequence is issued (and provided that sufficient substitution parameters exist), there is a probability that another short read sequence is issued. This probability decreases for each new sequence issued. Since the same random number generator seed is used across executions, the workload is deterministic.

The specified frequencies, implicitly define the query ratios between queries of different types, as well as a default target throughput. However the Test Sponsor may specify a different target throughput to test, by “squeezing” together or “stretching” apart the queries of the workload. This is achieved by means of the “Time Compression Ratio” that is multiplied by the frequencies (see Table 3.1). Therefore, different throughputs can be tested while maintaining the relative ratios between the different query types.

4 INTERACTIVE WORKLOAD

this empty space looks so bad, we should write something here – szarnyasg

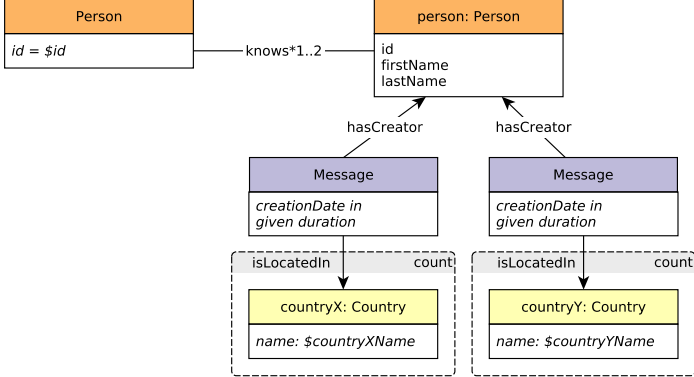
4.1 Complex Reads Query Descriptions

Query cards for complex read follow.

Interactive / complex / 1

query	Interactive / complex / 1			
title	Friends with certain name			
pattern	<pre> graph LR P1[Person id = \$id] -- knows*1..3 --> P2[person: Person firstName = \$firstName id lastName birthday creationDate gender browserUsed locationIP email speaks] P2 -- isLocatedIn --> Pl[Place name] P2 -- workAt --> Co[Company name] P2 -- studyAt --> Un[University name] Pl -- isLocatedIn --> Cn[Country name] Un -- isLocatedIn --> Ci[City name] </pre>			
desc.	Given a start Person, find Persons with a given first name that the start Person is connected to (excluding start Person) by at most 3 steps via Knows relationships. Return Persons, including summaries of the Persons workplaces and places of study.			
params	1	Person.id	ID	
	2	Person.firstName	String	
result	1	Person.id	ID	R
	2	Person.lastName	String	R
	3	Person.birthday	Date	R
	4	Person.creationDate	DateTime	R
	5	Person.gender	String	R
	6	Person.browserUsed	String	R
	7	Person.locationIP	String	R
	8	{Person.email}	{String}	R
	9	{Person.speaks}	{String}	R
	10	Person-isLocatedIn->Place.name	String	R
	11	{Person-studyAt->University.name, Person-studyAt->.classYear, Person-studyAt->University-isLocatedIn->City.name}	{}	R
	12	{Person-workAt->Company.name, Person-workAt->.workFrom, Person-workAt->Company-isLocatedIn->Country.name}	{}	R
sort	1	distanceFromPerson	↑	
	2	Person.lastName	↑	
	3	Person.id	↑	
limit	20			
CPs	1.3, 2.1, 5.3			
relevance	<p>This query is a representative of a simple navigational query. It looks for paths of length one, two or three through the knows relation, starting from a given Person and ending at a Person with a given name. It is interesting for several aspects. First, it requires for a complex aggregation, that is, returning the concatenation of universities, companies, languages and email information of the person. Second, it tests the ability of the optimizer to move the evaluation of sub-queries functionally dependant on the Person, after the evaluation of the top-k. Finally, performance is highly sensitive to properly estimating the cardinalities in each transitive path, and paying attention not to explode already visited Persons.</p>			

Interactive / complex / 3

query	Interactive / complex / 3			
title	Friends and friends of friends that have been to countries X and Y			
pattern	 <pre> graph TD P1[Person id = \$id] -- knows*1..2 --> P2[person: Person id firstName lastName] P2 -- hasCreator --> M1[Message creationDate in given duration] P2 -- hasCreator --> M2[Message creationDate in given duration] M1 -- isLocatedIn --> C1[countryX: Country name: \$countryXName] M2 -- isLocatedIn --> C2[countryY: Country name: \$countryYName] C1 --- countX[count] C2 --- countY[count] </pre>			
desc.	<p>Given a start Person, find Persons that are their friends and friends of friends (excluding start Person) that have made Posts/Comments in both of the given Countries, X and Y, within a given period. Only Persons that are foreign to Countries X and Y are considered, that is Persons whose Location is not Country X or Country Y.</p>			
params	1	Person.id	ID	
	2	CountryX.name	String	
	3	CountryY.name	String	
	4	startDate	Date	Beginning of requested period
	5	duration	32-bit Integer	Duration of requested period, in days the interval [startDate, startDate + Duration) is closed-open
result	1	Person.id	ID	R
	2	Person.firstName	String	R
	3	Person.lastName	String	R
	4	countX	32-bit Integer	A Number of Messages from Country X made by Person within the given time
	5	countY	32-bit Integer	A Number of Messages from Country Y made by Person within the given time
	6	count	32-bit Integer	A countX + countY
sort	1	countX	↓	
	2	Person.id	↑	
limit	20			
CPs	2.1, 3.1, 5.1			
relevance	<p>This query looks for paths of length two and three, starting from a Person, going to friends or friends of friends, and then moving to Messages. This query tests the ability of the query optimizer to select the most efficient join ordering, which will depend on the cardinalities of the intermediate results. Many friends of friends can be duplicate, then it is expected to eliminate duplicates and those people prior to access the Post and Comments, as well as eliminate those friends from countries X and Y, as the size of the intermediate results can be severely affected. A possible structural optimization could be to materialize the number of Posts and Comments created by a person, and progressively filter those people that could not even fall in the top 20 even having all their posts in the countries X and Y.</p>			

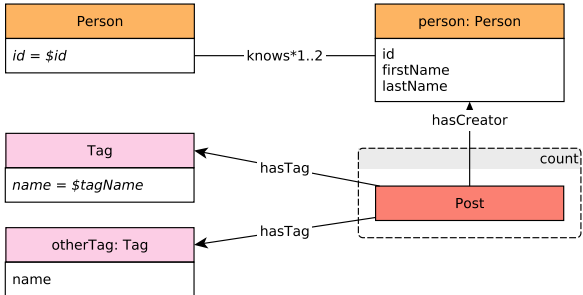
Interactive / complex / 4

query	Interactive / complex / 4			
title	New topics			
pattern	<pre> graph TD Person[id = \$id] -- knows --> person_Person[person: Person] person_Person -- hasCreator --> Post1[Post] person_Person -- hasCreator --> Post2[Post] Post1 -- hasTag --> Tag[Tag] Tag -.-> hasTag TODO: other friends' Post2 </pre>			
desc.	Given a start Person, find Tags that are attached to Posts that were created by that Person's friends. Only include Tags that were attached to friends' Posts created within a given time interval, and that were never attached to friends' Posts created before this interval.			
params	1	Person.id	ID	
	2	startDate	Date	
	3	duration	32-bit Integer	Duration of requested period, in days the interval [startDate, startDate + Duration) is closed-open
result	1	Tag.name	String	R
	2	count	32-bit Integer	A
	Number of Posts made within the given time interval that have this Tag			
sort	1	count	↓	
	2	Tag.name	↑	
limit	10			
CPs	2.3			
relevance	This query looks for paths of length two, starting from a given Person, moving to Posts and then to Tags. It tests the ability of the query optimizer to properly select the usage of hash joins or index based joins, depending on the cardinality of the intermediate results. These cardinalities are clearly affected by the input Person, the number of friends, the variety of Tags, the time interval and the number of Posts.			

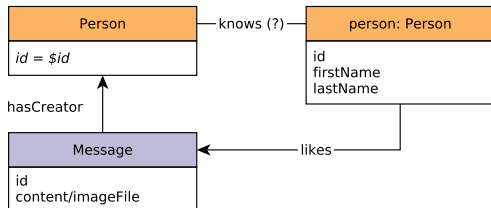
Interactive / complex / 5

query	Interactive / complex / 5				
title	New groups				
pattern					
desc.	<p>Given a start Person, find the Forums which that Person's friends and friends of friends (excluding start Person) became Members of after a given date. For each forum find the number of Posts that were created by any of these Persons. For each Forum and consider only those Persons which joined that particular Forum after the given date.</p>				
params	1	Person.id	ID		
	2	date	Date		
result	1	Forum.title	String	R	
	2	count	32-bit Integer	A	Number of Posts made in Forum that were created by friends
sort	1	count	↓		
	2	Forum.id	↑		
limit	20				
CPs	2.3, 3.3				
relevance	<p>This query looks for paths of length two and three, starting from a given Person, moving to friends and friends of friends, and then getting the Forums they are members of. Besides testing the ability of the query optimizer to select the proper join operator, it rewards the usage of indexes, but their accesses will be presumably scattered due to the two/three-hop search space of the query, leading to unpredictable and scattered index accesses. Having efficient implementations of such indexes will be highly beneficial.</p>				

Interactive / complex / 6

query	Interactive / complex / 6													
title	Tag co-occurrence													
pattern														
desc.	Given a start Person and some Tag, find the other Tags that occur together with this Tag on Posts that were created by start Person’s friends and friends of friends (excluding start Person). Return For each Tag, find the count of Posts that were created by these Persons, which contain both this Tag and the given Tag.													
params	<table><tr><td>1</td><td>Person.id</td><td>ID</td><td></td></tr><tr><td>2</td><td>Tag.name</td><td>String</td><td></td></tr></table>	1	Person.id	ID		2	Tag.name	String						
1	Person.id	ID												
2	Tag.name	String												
result	<table><tr><td>1</td><td>Tag.name</td><td>String</td><td>R</td><td></td></tr><tr><td>2</td><td>count</td><td>32-bit Integer</td><td>A</td><td>Number of Posts that were created by friends and friends of friends, which contain this Tag</td></tr></table>	1	Tag.name	String	R		2	count	32-bit Integer	A	Number of Posts that were created by friends and friends of friends, which contain this Tag			
1	Tag.name	String	R											
2	count	32-bit Integer	A	Number of Posts that were created by friends and friends of friends, which contain this Tag										
sort	<table><tr><td>1</td><td>count</td><td>↓</td><td></td></tr><tr><td>2</td><td>Tag.name</td><td>↑</td><td></td></tr></table>	1	count	↓		2	Tag.name	↑						
1	count	↓												
2	Tag.name	↑												
limit	10													
CPs	5.1													
relevance	This query looks for paths of lengths three or four, starting from a Given Person, moving to friends or friends of friends, then to Posts and finally ending at a given Tag.													

Interactive / complex / 7

query	Interactive / complex / 7																																												
title	Recent likers																																												
pattern	 <pre>graph TD Person1[Person id = \$id] -- knows (?) --- person2[person: Person id firstName lastName] Message[Message id content/imageFile] -- hasCreator --> Person1 person2 -- likes --> Message</pre>																																												
desc.	Given a start Person, find (most recent) Likes on any of start Person’s Messages. Find Persons that Liked any of start Person’s Messages, the Messages they liked most recently, creation date of that Like, and the latency (in minutes) between creation of Messages and Like. Additionally, for each Person found return a flag indicating whether the liker is a friend of start Person. In the case that a Person Liked multiple Messages at the same time, return the Message with lowest identifier.																																												
params	<table><tr><td>1</td><td>Person.id</td><td>64-bit Integer</td><td></td></tr></table>					1	Person.id	64-bit Integer																																					
1	Person.id	64-bit Integer																																											
result	<table><tr><td>1</td><td>Person.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>2</td><td>Person.firstName</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>Person.lastName</td><td>String</td><td>R</td><td></td></tr><tr><td>4</td><td>Like.creationDate</td><td>DateTime</td><td>R</td><td></td></tr><tr><td>5</td><td>Message.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>6</td><td>Message.content or Post.imageFile</td><td>String</td><td>R</td><td></td></tr><tr><td>7</td><td>latency</td><td>32-bit Integer</td><td>C</td><td>Duration between creation of Message and Like, in minutes</td></tr><tr><td>8</td><td>isNew</td><td>Boolean</td><td>C</td><td>false if liker Person is friend of start Person, true otherwise</td></tr></table>					1	Person.id	ID	R		2	Person.firstName	String	R		3	Person.lastName	String	R		4	Like.creationDate	DateTime	R		5	Message.id	ID	R		6	Message.content or Post.imageFile	String	R		7	latency	32-bit Integer	C	Duration between creation of Message and Like, in minutes	8	isNew	Boolean	C	false if liker Person is friend of start Person, true otherwise
1	Person.id	ID	R																																										
2	Person.firstName	String	R																																										
3	Person.lastName	String	R																																										
4	Like.creationDate	DateTime	R																																										
5	Message.id	ID	R																																										
6	Message.content or Post.imageFile	String	R																																										
7	latency	32-bit Integer	C	Duration between creation of Message and Like, in minutes																																									
8	isNew	Boolean	C	false if liker Person is friend of start Person, true otherwise																																									
sort	<table><tr><td>1</td><td>Like.creationDate</td><td>↓</td><td></td></tr><tr><td>2</td><td>Person.id</td><td>↑</td><td></td></tr></table>					1	Like.creationDate	↓		2	Person.id	↑																																	
1	Like.creationDate	↓																																											
2	Person.id	↑																																											
limit	20																																												
CPs	2.2, 2.3, 3.3, 5.1																																												
relevance	This query looks for paths of length two, starting from a given Person, moving to its published messages and then to Persons who liked them. It tests several aspects related to join optimization, both at query optimization plan level and execution engine level. On the one hand, many of the columns needed for the projection are only needed in the last stages of the query, so the optimizer is expected to delay the projection until the end. This query implies accessing 2-hop data, and as a consequence, index accesses are expected to be scattered. We expect to observe variate cardinalities, depending on the characteristics of the input parameter, so properly selecting the join operators will be crucial. This query has a lot of correlated sub-queries, so it is testing the ability to flatten the query execution plans.																																												

Interactive / complex / 8

query	Interactive / complex / 8			
title	Recent replies			
pattern	<pre> graph TD P1[Person id = \$id] -- hasCreator --> M[Message] M -- replyOf --> C[comment: Comment id content creationDate] C -- hasCreator --> P2[person: Person id firstName lastName] </pre>			
desc.	Given a start Person, find (most recent) Comments that are replies to Messages of the start Person. Only consider immediate (1-hop) replies, not the transitive (multi-hop) case. Return the reply Comments, and the Person that created each reply Comment.			
params	1	Person.id	ID	
result	1	Person.id	ID	R
	2	Person.firstName	String	R
	3	Person.lastName	String	R
	4	Comment.creationDate	DateTime	R
	5	Comment.id	ID	R
	6	Comment.content	String	R
sort	1	Comment.creationDate	↓	
	2	Comment.id	↑	
limit	20			
CPs	2.4, 3.2, 3.3, 5.3			
relevance	This query looks for paths of length two, starting from a given Person, going through its created Messages and finishing at their replies. In this query there is temporal locality between the replies being accessed. Thus the top k order by this can interact with the selection, i.e. do not consider older posts than the 20th oldest seen so far.			

Interactive / complex / 9

query	Interactive / complex / 9			
title	Recent posts and comments by friends or friends of friends			
pattern	<pre> graph LR P[Person id = \$id] -- knows*1..2 --> PP[person: Person id firstName lastName] PP -- hasCreator --> M[Message id content / imageFile creationDate] </pre>			
desc.	Given a start Person, find the (most recent) Messages created by that Person's friends or friends of friends (excluding start Person). Only consider the Messages created before a given date (excluding that date).			
params	1	Person.id	ID	
	2	date	Date	
result	1	Message-hasCreator->Person.id	ID	R
	2	Message-hasCreator->Person.firstName	String	R
	3	Message-hasCreator->Person.lastName	String	R
	4	Message.id	ID	R
	5	Message.content or Post.imageFile	String	R
	6	Message.creationDate	DateTime	R
sort	1	Message.creationDate	↓	
	2	Message.id	↑	
limit	20			
CPs	1.1, 1.2, 2.2, 2.3, 3.2, 3.3			
relevance	This query looks for paths of length two or three, starting from a given Person, moving to its friends and friends of friends, and ending at their created Messages. This is one of the most complex queries, as the list of choke-points indicates. This query is expected to touch variable amounts of data with entities of different characteristics, and therefore, properly estimating cardinalities and selecting the proper operators will be crucial.			

Interactive / complex / 10

query	Interactive / complex / 10			
title	Friend recommendation			
pattern	<div><div><div>Person</div><div>$id = \\$id$</div></div><div><div>person: Person</div><div>$birthDay \text{ cond's}$ id firstName lastName gender</div></div><div>knows*2..2</div></div> <div><div><div>Person</div><div><div>hasCreator</div><div>Post</div><div>hasTag</div><div>Tag</div><div>count</div></div></div><div>Person</div><div>hasInterest</div></div>			
desc.	<p>Given a start Person, find that Person’s friends of friends (excluding start Person, and immediate friends), who were born on or after the 21st of a given month (in any year) and before the 22nd of the following month. Calculate the similarity between each of these Persons and start Person, where similarity for any Person is defined as follows:</p> <ul style="list-style-type: none">• common = number of Posts created by that Person, such that the Post has a Tag that start Person is Interested in• uncommon = number of Posts created by that Person, such that the Post has no Tag that start Person is Interested in• similarity = common - uncommon			
params	<div><div>1</div><div>Person.id</div><div>ID</div></div> <div><div>2</div><div>month</div><div>32-bit Integer</div></div>	<div></div> <div>Between 1-12</div>		
result	<div><div>1</div><div>Person.id</div><div>ID</div></div> <div><div>2</div><div>Person.firstName</div><div>String</div></div> <div><div>3</div><div>Person.lastName</div><div>String</div></div> <div><div>4</div><div>similarity</div><div>32-bit Integer</div></div> <div><div>5</div><div>Person.gender</div><div>String</div></div> <div><div>6</div><div>Person-isLocatedIn->Place.name</div><div>String</div></div>	<div>R</div> <div>R</div> <div>R</div> <div>C</div> <div>R</div> <div>R</div>		
sort	<div><div>1</div><div>similarity</div><div>↓</div></div> <div><div>2</div><div>Person.id</div><div>↑</div></div>			
limit	10			
CPs	2.3, 3.3, 4.1, 4.2, 5.1, 5.2, 6.1, 7.1			
relevance	<p>This query looks for paths of length two, starting from a Person and ending at the friends of their friends. It does widely scattered graph traversal, and one expects no locality of in friends of friends, as these have been acquired over a long time and have widely scattered identifiers. The join order is simple but one must see that the anti-join for "not in my friends" is better with hash. Also the last pattern in the scalar sub-queries joining or anti-joining the tags of the candidate’s posts to interests of self should be by hash.</p>			

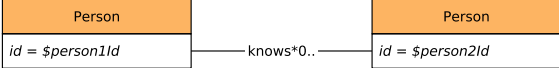
Interactive / complex / 11

query	Interactive / complex / 11			
title	Job referral			
pattern	<pre> graph LR Person[id = \$id] -- knows*1..2 --> person["person: Person id firstName lastName"] person -- "worksAt year.worksFrom <= \$year" --> Organisation[name] Organisation -- isLocatedIn --> Country[name = \$name] </pre>			
desc.	Given a start Person, find that Person's friends and friends of friends (excluding start Person) who started Working in some Company in a given Country, before a given date (year).			
params	1	Person.id	ID	
	2	Country.name	String	
	3	year	32-bit Integer	
result	1	Person.id	ID	R
	2	Person.firstName	String	R
	3	Person.lastName	String	R
	4	Person-worksAt->Organisation.name	String	R
	5	Person-worksAt->.worksFrom	32-bit Integer	R
sort	1	Person-worksAt->.worksFrom	↑	
	2	Person.id	↑	
	3	Person-worksAt->Organisation.name	↓	
limit	10			
CPs	1.4, 2.3, 2.4, 3.3			
relevance	This query looks for paths of length two or three, starting from a Person, moving to friends or friends of friends, and ending at a Company. In this query, there are selective joins and a top k order by that can be exploited for optimizations.			

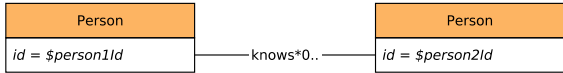
Interactive / complex / 12

query	Interactive / complex / 12				
title	Expert search				
pattern					
desc.	<p>Given a start Person, find the Comments that this Person's friends made in reply to Posts, considering only those Comments that are immediate (1-hop) replies to Posts, not the transitive (multi-hop) case. Only consider Posts with a Tag in a given TagClass or in a descendant of that TagClass. Count the number of these reply Comments, and collect the Tags that were attached to the Posts they replied to, but only collect Tags with the given TagClass or with a descendant of that TagClass. Return Persons with at least one reply, the reply count, and the collection of Tags.</p>				
params	1	Person.id	ID		
	2	TagClass.name	String		
result	1	Person.id	ID	R	
	2	Person.firstName	String	R	
	3	Person.lastName	String	R	
	4	{Tag.name}	{String}	R	
	5	count	32-bit Integer	A	Number of reply Comments
sort	1	count	↓		
	2	Person.id	↑		
limit	20				
CPs	1.5, 3.3, 7.2, 7.3				
relevance	<p>This query looks for paths of length three, starting at a Person, moving to its friends, the to their Comments and ending at the Post the Comments are replying. The chain from original post to the reply is transitive. The traversal may be initiated at either end, the system may note that this is a tree, hence leaf to root is always best. Additionally, a hash table can be built from either end, e.g. from the friends of self, from the tags in the category, from the or other.</p>				

Interactive / complex / 13

query	Interactive / complex / 13			
title	Single shortest path			
pattern				
desc.	<p>Given two Persons, find the shortest path between these two Persons in the subgraph induced by the Knows relationships. Return the length of this path:</p> <ul style="list-style-type: none"> • -1 : no path found • 0: start person = end person • > 0: regular case 			
params	1	person1.id	ID	
	2	person2.id	ID	
result	1	length	32-bit Integer	C
CPs	3.3, 7.2, 7.3			
relevance	This query looks for a variable length path, starting at a given Person and finishing at an another given Person. Proper cardinality estimation and search space pruning, will be crucial. This query also allows for possible parallel implementations.			

Interactive / complex / 14

query	Interactive / complex / 14			
title	Weighted/unweighted paths			
pattern				
desc.	<p>Given two Persons, find all (unweighted) shortest paths between these two Persons, in the sub-graph induced by the Knows relationship. Then, for each path calculate a weight. The nodes in the path are Persons, and the weight of a path is the sum of weights between every pair of consecutive Person nodes in the path. The weight for a pair of Persons is calculated such that every reply (by one of the Persons) to a Post (by the other Person) contributes 1.0, and every reply (by ones of the Persons) to a Comment (by the other Person) contributes 0.5. Return all the paths with shortest length, and their weights. Do not return any rows if there is now path between the two Persons.</p>			
params	1	person1.id	ID	
	2	person2.id	ID	
result	1	[Person.id]	[ID]	C Identifiers representing an ordered sequence of the Persons in the path
	2	weight	64-bit Float	C
sort	1	weight	↓	The order of paths with the same weight is unspecified
CPs	3.3, 7.2, 7.3			
relevance	<p>This query looks for a variable length path, starting at a given Person and finishing at an another given Person. This is a more complex query as not only requires computing the path length, but returning it and computing a weight. To compute this weight one must look for smaller sub-queries with paths of length three, formed by the two Persons at each step, a Post and a Comment.</p>			

4.2 Short Reads Query Descriptions

Interactive / short / 1

query	Interactive / short / 1				
title	Person Profile				
pattern	<div><div>person: Person</div><div>id lastName birthday ipAddress browser</div><div>→ isLocatedIn →</div><div>city: City</div><div>name</div></div>				
desc.	Given a start Person, retrieve their first name, last name, birthday, IP address, browser, and city of residence.				
params	<div><div>1</div><div>Person.id</div><div>ID</div><div></div></div>				
result	1	Person.firstName	String	R	
	2	Person.lastName	String	R	
	3	Person.birthday	Date	R	
	4	Person.locationIP	String	R	
	5	Person.browserUsed	String	R	
	6	Person-isLocatedIn->Place.id	32-bit Integer	R	
	7	Person.gender	String	R	
	8	Person.creationDate	DateTime	R	

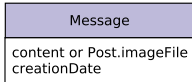
Interactive / short / 2

query	Interactive / short / 2																																							
title	Person Recent Messages																																							
pattern	<pre>graph TD person[person: Person] -- hasCreator --> Message[Message] Message -- replyOf*1.. --> Post[Post] Post -- hasCreator --> op[op: Person]</pre> <p>The diagram illustrates the query pattern. It shows three entity types: person: Person (orange box), Message (purple box), and Post (red box). The person: Person entity has attributes: id, lastName, birthday, ipAddress, and browser. The Message entity has attributes: id, lastName, birthday, ipAddress, and browser. The Post entity has no attributes listed. Relationships are shown as follows: person: Person has a hasCreator relationship with Message. Message has a replyOf*1.. relationship with Post. Post has a hasCreator relationship with op: Person (orange box).</p>																																							
desc.	Given a start Person, retrieve the last 10 Messages created by that user. For each message, return that message, the original post in its conversation, and the author of that post. If any of the Messages is a Post, then the original Post will be the same Message, i.e. that Message will appear twice in that result.																																							
params	<table><tr><td>1</td><td>Person.id</td><td>ID</td><td></td></tr></table>					1	Person.id	ID																																
1	Person.id	ID																																						
result	<table><tr><td>1</td><td>Message.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>Message.content or Post.imageFile</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>Message.creationDate</td><td>DateTime</td><td>R</td><td></td></tr><tr><td>4</td><td>Post.id or Comment-replyOf*->Post.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>5</td><td>Post-hasCreator->Person.id or Comment-replyOf*->Post-hasCreator->Person.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>6</td><td>Post-hasCreator->Person.firstName or Comment-replyOf*->Post-hasCreator->Person.firstName</td><td>String</td><td>R</td><td></td></tr><tr><td>7</td><td>Post-hasCreator->Person.lastName or Comment-replyOf*->Post-hasCreator->Person.lastName</td><td>String</td><td>R</td><td></td></tr></table>					1	Message.id	64-bit Integer	R		2	Message.content or Post.imageFile	String	R		3	Message.creationDate	DateTime	R		4	Post.id or Comment-replyOf*->Post.id	ID	R		5	Post-hasCreator->Person.id or Comment-replyOf*->Post-hasCreator->Person.id	ID	R		6	Post-hasCreator->Person.firstName or Comment-replyOf*->Post-hasCreator->Person.firstName	String	R		7	Post-hasCreator->Person.lastName or Comment-replyOf*->Post-hasCreator->Person.lastName	String	R	
1	Message.id	64-bit Integer	R																																					
2	Message.content or Post.imageFile	String	R																																					
3	Message.creationDate	DateTime	R																																					
4	Post.id or Comment-replyOf*->Post.id	ID	R																																					
5	Post-hasCreator->Person.id or Comment-replyOf*->Post-hasCreator->Person.id	ID	R																																					
6	Post-hasCreator->Person.firstName or Comment-replyOf*->Post-hasCreator->Person.firstName	String	R																																					
7	Post-hasCreator->Person.lastName or Comment-replyOf*->Post-hasCreator->Person.lastName	String	R																																					
sort	<table><tr><td>1</td><td>Message.creationDate</td><td>↓</td><td></td></tr><tr><td>2</td><td>Message.id</td><td>↓</td><td></td></tr></table>					1	Message.creationDate	↓		2	Message.id	↓																												
1	Message.creationDate	↓																																						
2	Message.id	↓																																						

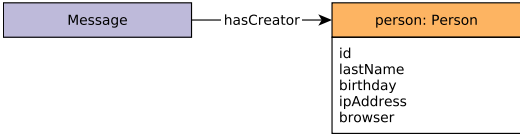
Interactive / short / 3

query	Interactive / short / 3																								
title	Person Friends																								
pattern	<div><div><div>Person</div><div><i>id = \$id</i></div></div><div>— knows — creationDate</div><div><div>person: Person</div><div>id firstName lastName</div></div></div>																								
desc.	Given a start Person, retrieve all of their friends, and the date at which they became friends.																								
params	<div><div>1</div><div>Person.id</div><div>ID</div><div></div></div>																								
result	<table><tr><td>1</td><td>Person.id</td><td>ID</td><td>R</td><td></td></tr><tr><td>2</td><td>Person.firstName</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>Person.lastName</td><td>String</td><td>R</td><td></td></tr><tr><td>4</td><td>Knows.creationDate</td><td>DateTime</td><td>R</td><td></td></tr></table>					1	Person.id	ID	R		2	Person.firstName	String	R		3	Person.lastName	String	R		4	Knows.creationDate	DateTime	R	
1	Person.id	ID	R																						
2	Person.firstName	String	R																						
3	Person.lastName	String	R																						
4	Knows.creationDate	DateTime	R																						
sort	<table><tr><td>1</td><td>Knows.creationDate</td><td>↓</td><td></td></tr><tr><td>2</td><td>Person.id</td><td>↑</td><td></td></tr></table>					1	Knows.creationDate	↓		2	Person.id	↑													
1	Knows.creationDate	↓																							
2	Person.id	↑																							

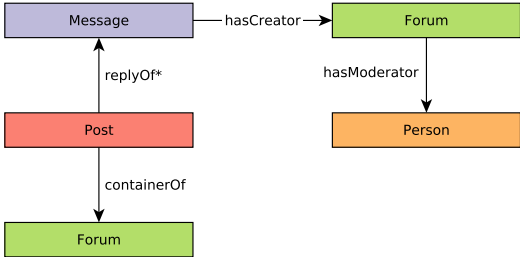
Interactive / short / 4

query	Interactive / short / 4													
title	Message Content													
pattern														
desc.	Given a Message, retrieve its content and creation date.													
params	<div>1</div> <div>Message.id</div> <div>ID</div> <div></div>													
result	<table> <tr> <td>1</td><td>Message.creationDate</td><td>ID</td><td>R</td><td></td></tr> <tr> <td>2</td><td>Message.content or Post.imageFile</td><td>String</td><td>R</td><td></td></tr> </table>				1	Message.creationDate	ID	R		2	Message.content or Post.imageFile	String	R	
1	Message.creationDate	ID	R											
2	Message.content or Post.imageFile	String	R											

Interactive / short / 5

query	Interactive / short / 5				
title	Message Creator				
pattern					
desc.	Given a Message, retrieve its author.				
params	<div> <div>1</div> <div>Message.id</div> <div>ID</div> <div></div> </div>				
result	1	Message-hasCreator->Person.id	ID	R	
	2	Message-hasCreator->Person.firstName	String	R	
	3	Message-hasCreator->Person.lastName	String	R	

Interactive / short / 6

query	Interactive / short / 6				
title	Message Forum				
pattern					
desc.	Given a Message, retrieve the Forum that contains it and the Person that moderates that forum. Since comments are not directly contained in forums, for comments, return the forum containing the original post in the thread which the comment is replying to.				
params	<div> <div>1</div> <div>Message.id</div> <div>ID</div> <div></div> </div>				
result	1	Message<-containerOf-Forum.id	ID	R	
	2	Message<-containerOf-Forum.title	String	R	
	3	Message<-containerOf-Forum-hasModerator->Person.id	ID	R	
	4	Message<-containerOf-Forum-hasModerator->Person.firstName	String	R	
	5	Message<-containerOf-Forum-hasModerator->Person.lastName	String	R	

Interactive / short / 7

query	Interactive / short / 7				
title	Message Replies				
pattern	<pre>graph TD p1[p1: Person] -- hasCreator --> Message[Message] p2[p2: Person] -- hasCreator --> Comment[Comment] Message -- replyOf --> Comment p1 -- knows --> p2</pre>				
desc.	Given a Message, retrieve the (1-hop) Comments that reply to it. In addition, return a boolean flag knows indicating if the author of the reply knows the author of the original message. If author is same as original author, return false for knows flag.				
params	1	Message.id	ID		
result	1	Message<-replyOf-Comment.id	ID	R	
	2	Message<-replyOf-Comment.content	String	R	
	3	Message<-replyOf-Comment.creationDate	DateTime	R	
	4	Comment-hasCreator->Person.id	ID	R	
	5	Comment-hasCreator->Person.firstName	String	R	
	6	Comment-hasCreator->Person.lastName	String	R	
	7	knows	Boolean	C	Original message author knows reply author
sort	1	Message<-replyOf-Comment.creationDate	↓		
	2	Message-hasCreator->Person.id	↑		

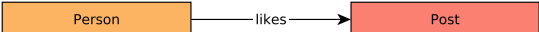
4.3 Update Query Descriptions

Each update query inserts either (1) a single vertex of a certain type, along with its edges to other existing vertices or (2) a single edge of a certain type between two existing vertices.


Interactive / update / 1

query	Interactive / update / 1			
title	Add Person			
pattern	<div><div><div>Person</div><div>id firstName lastName gender birthDay creationDate locationIp browserUsed speaks email</div></div><div><div>isLocatedIn</div><div>hasInterest</div><div>worksAt</div><div>studyAt</div></div><div><div>City</div><div>Tag</div><div>Company</div><div>University</div></div></div>			
desc.	Add a Person to the social network.			
params	1	Person.id	ID	
	2	Person.firstName	String	
	3	Person.lastName	String	
	4	Person.gender	String	
	5	Person.birthDay	Date	
	6	Person.creationDate	DateTime	
	7	Person.locationIp	String	
	8	Person.browserUsed	String	
	9	Person-isLocatedIn->City.id	ID	
	10	Person.speaks	{String}	
	11	Person.email	{String}	
	12	Person-hasInterest->Tag.id	{ID}	
	13	{Person-studyAt->University.id, Person-studyAt->.classYear}	{ID, 32-bit Integer}	
	14	{Person-workAt->Company.id, Person-workAt->.workFrom}	{ID, 32-bit Integer}	

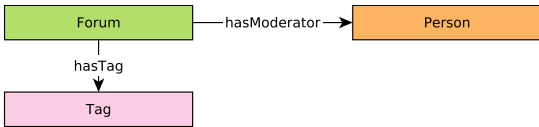
Interactive / update / 2

query	Interactive / update / 2			
title	Add Post Like			
pattern				
desc.	Add a Like to a Post of the social network.			
params	1	Person.id	ID	
	2	Post.id	ID	
	3	Person-likes->.creationDate	DateTime	

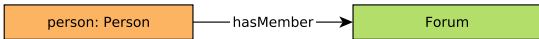
Interactive / update / 3

query	Interactive / update / 3		
title	Add Comment Like		
pattern			
desc.	Add a Like to a Comment of the social network.		
params	1	Person.id	ID
	2	Comment.id	ID
	3	Person-likes->.creationDate	DateTime

Interactive / update / 4

query	Interactive / update / 4		
title	Add Forum		
pattern			
desc.	Add a Forum to the social network.		
params	1	Forum.id	ID
	2	Forum.title	String
	3	Forum.creationDate	DateTime
	4	Forum-hasModerator->Person.id	ID
	5	Forum-hasTag->Tag.id	{ID}

Interactive / update / 5

query	Interactive / update / 5		
title	Add Forum Membership		
pattern			
desc.	Add a Forum membership to the social network.		
params	1	Person.id	ID
	2	Forum-hasMember->Person.id	ID
	3	Forum-hasMember->.joinDate	DateTime

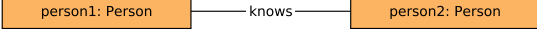
Interactive / update / 6

query	Interactive / update / 6				
title	Add Post				
pattern	<div><div><div>Post</div><div>id imageFile creationDate locationIp browserUsed language content length</div></div><div><div>Forum</div><div>containerOf</div></div><div><div>Person</div><div>hasCreator</div></div><div><div>Country</div><div>isLocatedIn</div></div><div><div>Tag</div><div>hasTag</div></div></div>				
desc.	Add a Post to the social network.				
params	1	Post.id	ID		
	2	Post.imageFile	String		
	3	Post.creationDate	DateTime		
	4	Post.locationIp	String		
	5	Post.browserUsed	String		
	6	Post.language	String		
	7	Post.content	Text		
	8	Post.length	32-bit Integer		
	9	Post-hasCreator->Person.id	ID		
	10	Post<-containerOf-Forum.id	ID		
	11	Post-isLocatedIn->Country.id	ID		
	12	{Post-hasTag->Tag.id}	{ID}		

Interactive / update / 7

query	Interactive / update / 7			
title	Add Comment			
pattern	<div><div><div>Comment</div><div>id imageFile creationDate locationIp browserUsed language content length</div></div><div><div>Forum</div><div>containerOf</div></div><div><div>person: Person</div><div>hasCreator</div></div><div><div>Country</div><div>isLocatedIn</div></div><div><div>Tag</div><div>hasTag</div></div></div>			
desc.	Add a Comment replying to a Post/Comment to the social network.			
params	1	Comment.id	ID	
	2	Comment.creationDate	DateTime	
	3	Comment.locationIp	String	
	4	Comment.browserUsed	String	
	5	Comment.content	Text	
	6	Comment.length	32-bit Integer	
	7	Comment-hasCreator->Person.id	ID	
	8	Comment-isLocatedIn->Country.id	ID	
	9	Comment-replyOf->Post.id	ID	-1 if the comment is a reply of a comment
	10	Comment-replyOf->Comment.id	ID	-1 if the comment is a reply of a post
	11	{Comment-hasTag->Tag.id}	{ID}	

Interactive / update / 8

query	Interactive / update / 8			
title	Add Friendship			
pattern				
desc.	Add a friendship relation to the social network			
params	1	Person.id	ID	Person 1
	2	Person.id	ID	Person 2
	3	Person-knows->.creationDate	DateTime	

5 BUSINESS INTELLIGENCE WORKLOAD

5.1 Read Query Descriptions

BI / read / 1

query	BI / read / 1								
title	Posting summary								
pattern	<table><tr><td>message: Message</td></tr><tr><td>creationDate < \$date</td></tr><tr><td>length year(creationDate)</td></tr></table>					message: Message	creationDate < \$date	length year(creationDate)	
message: Message									
creationDate < \$date									
length year(creationDate)									
desc.	<p>Given a date, find all <i>Messages</i> created before that date. Group them by a 3-level grouping:</p> <ul style="list-style-type: none">1. by year of creation2. for each year, group into <i>Message</i> types (<i>Posts</i> or <i>Comments</i>)3. for each year-type group, split into four groups based on length of their content<ul style="list-style-type: none">1. 0 <= length < 40: short2. 40 <= length < 80: one liner3. 80 <= length < 160: tweet4. 160 <= length: long								
params	<table><tr><td>1</td><td>date</td><td>Date</td><td></td></tr></table>					1	date	Date	
1	date	Date							
result	1	messageYear	32-bit Integer	R	Year of the <i>Message</i>				
	2	messageType	String	M	post/comment (in lowercase)				
	3	lengthCategory	String	C	short/one-liner/tweet/long (in lowercase)				
	4	messageCount	32-bit Integer	A	Total number of <i>Messages</i> (<i>Posts/Comments</i>) in that group				
	5	averageMessageLength	32-bit Integer	A	Average length of the <i>Message</i> content in that group				
	6	sumMessageLength	32-bit Integer	A	Sum of all <i>Message</i> content lengths				
	7	percentageOfMessages	32-bit Float	A	Number of <i>Messages</i> in group as a percentage of all messages created before the given date				
sort	1	year	↓						
	2	messageType	↑	first <i>Posts</i> , then <i>Comments</i>					
	3	lengthCategory	↑	order based on the length of the category, short, one liner, etc.					
CPs	1.2, 3.2, 4.1								

BI / read / 2

query	BI / read / 2																																		
title	Top tags for country, age, gender, time																																		
pattern	<div><div><div>country: Country</div><div><div>name = \$country1 or name = \$country2</div><div>name</div></div></div><div><div>City</div><div>isPartOf</div><div>country: Country</div></div><div><div>person: Person</div><div><div>gender birthday</div></div></div><div><div>isLocatedIn</div><div>City</div><div>person: Person</div></div><div><div>hasCreator</div><div>person: Person</div><div>message: Message</div></div><div><div>message: Message</div><div><div>creationDate >= \$date1 and creationDate <= \$date2</div><div>month(creationDate)</div></div></div><div><div>hasTag</div><div>message: Message</div><div>tag: Tag</div></div><div><div>messageCount = count</div><div>message: Message</div><div>tag: Tag</div></div><div><div>messageCount > 100</div><div>message: Message</div><div>tag: Tag</div></div></div>																																		
desc.	<p>Select all <i>Messages</i> (Posts & Comments) created between date1-date2 (inclusive) by <i>Persons</i> located in country1 or country2. Select the creator <i>Persons</i> and the <i>Tags</i> of these <i>Messages</i>. Split these <i>Persons</i>, <i>Tags</i> and <i>Messages</i> into a 5-level grouping:</p> <div><div>1. name of country of <i>Person</i>,</div><div>2. month the <i>Message</i> was created,</div><div>3. gender of <i>Person</i>,</div><div>4. age group of <i>Person</i>, defined as years between person’s birthday and end of simulation (2013-01-01), divided by 5, rounded down,</div><div>5. name of tag attached to <i>Message</i>.</div></div> <p>Consider only those groups where number of <i>Messages</i> is greater than 100.</p>																																		
params	<table><tr><td>1</td><td>date1</td><td>Date</td><td></td></tr><tr><td>2</td><td>date2</td><td>Date</td><td></td></tr><tr><td>3</td><td>country1</td><td>String</td><td></td></tr><tr><td>4</td><td>country2</td><td>String</td><td></td></tr></table>					1	date1	Date		2	date2	Date		3	country1	String		4	country2	String															
1	date1	Date																																	
2	date2	Date																																	
3	country1	String																																	
4	country2	String																																	
result	<table><tr><td>1</td><td>country.name</td><td>String</td><td>R</td><td></td></tr><tr><td>2</td><td>messageMonth</td><td>32-bit Integer</td><td>R</td><td>1–12</td></tr><tr><td>3</td><td>person.gender</td><td>String</td><td>R</td><td>male/female</td></tr><tr><td>4</td><td>ageGroup</td><td>32-bit Integer</td><td>C</td><td></td></tr><tr><td>5</td><td>tag.name</td><td>String</td><td>R</td><td></td></tr><tr><td>6</td><td>messageCount</td><td>64-bit Integer</td><td>A</td><td>The number of messages in the group</td></tr></table>					1	country.name	String	R		2	messageMonth	32-bit Integer	R	1–12	3	person.gender	String	R	male/female	4	ageGroup	32-bit Integer	C		5	tag.name	String	R		6	messageCount	64-bit Integer	A	The number of messages in the group
1	country.name	String	R																																
2	messageMonth	32-bit Integer	R	1–12																															
3	person.gender	String	R	male/female																															
4	ageGroup	32-bit Integer	C																																
5	tag.name	String	R																																
6	messageCount	64-bit Integer	A	The number of messages in the group																															
sort	<table><tr><td>1</td><td>messageCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>tag.name</td><td>↑</td><td></td></tr><tr><td>3</td><td>ageGroup</td><td>↑</td><td></td></tr><tr><td>4</td><td>person.gender</td><td>↑</td><td></td></tr><tr><td>5</td><td>messageMonth</td><td>↑</td><td></td></tr><tr><td>6</td><td>country.name</td><td>↑</td><td></td></tr></table>					1	messageCount	↓		2	tag.name	↑		3	ageGroup	↑		4	person.gender	↑		5	messageMonth	↑		6	country.name	↑							
1	messageCount	↓																																	
2	tag.name	↑																																	
3	ageGroup	↑																																	
4	person.gender	↑																																	
5	messageMonth	↑																																	
6	country.name	↑																																	
limit	100																																		
CPs	1.1, 1.2, 1.4, 2.1, 2.3, 3.1, 3.2																																		

BI / read / 3

query	BI / read / 3				
title	Tag evolution				
pattern	<div><div><div>countMonth1 = count</div><div><div>Message</div><div>year(creationDate) = \$year and month(creationDate) = \$month</div></div><div>hasTag</div><div><div>tag: Tag</div><div>name</div></div></div><div>OR</div><div><div>countMonth2 = count</div><div><div>Message</div><div>year(creationDate) = \$year + \$month/12 and month(creationDate) = \$month + \$month%12</div></div><div>hasTag</div><div><div>tag: Tag</div><div>name</div></div></div></div>				
desc.	Find the <i>Tags</i> that were used in <i>Messages</i> during the given month of the given year and the <i>Tags</i> that were used during the next month. For both months, compute the count of <i>Messages</i> that used each of the <i>Tags</i> .				
params	<div><div>1</div><div>year</div><div>32-bit Integer</div><div></div></div> <div><div>2</div><div>month</div><div>32-bit Integer</div><div></div></div>				
result	<div><div>1</div><div>tag.name</div><div>String</div><div>R</div><div></div></div> <div><div>2</div><div>countMonth1</div><div>32-bit Integer</div><div>A</div><div>Occurrences of the tag during the given year and month</div></div> <div><div>3</div><div>countMonth2</div><div>32-bit Integer</div><div>A</div><div>Occurrences of the tag during the next month after the given year and month</div></div> <div><div>4</div><div>diff</div><div>32-bit Integer</div><div>A</div><div>Absolute difference of countMonth1 and countMonth2</div></div>				
sort	<div><div>1</div><div>diff</div><div>↓</div><div></div></div> <div><div>2</div><div>tag.name</div><div>↑</div><div></div></div>				
limit	100				
CPs	2.4, 3.1, 3.2, 4.1, 4.3, 5.3, 6.1				

BI / read / 4

query	BI / read / 4																													
title	Popular topics in a country																													
pattern																														
desc.	Given a <i>TagClass</i> and a <i>Country</i> , find all the <i>Forums</i> created in the given <i>Country</i> , containing at least one <i>Post</i> with <i>Tags</i> belonging directly to the given <i>TagClass</i> . The location of a <i>Forum</i> is identified by the location of the <i>Forum</i> 's moderator.																													
params	<table><tr><td>1</td><td>tagClass</td><td>String</td><td></td></tr><tr><td>2</td><td>country</td><td>String</td><td></td></tr></table>					1	tagClass	String		2	country	String																		
1	tagClass	String																												
2	country	String																												
result	<table><tr><td>1</td><td>forum.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>forum.title</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>forum.creationDate</td><td>DateTime</td><td>R</td><td></td></tr><tr><td>4</td><td>person.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>5</td><td>postCount</td><td>32-bit Integer</td><td>A</td><td></td></tr></table>					1	forum.id	64-bit Integer	R		2	forum.title	String	R		3	forum.creationDate	DateTime	R		4	person.id	64-bit Integer	R		5	postCount	32-bit Integer	A	
1	forum.id	64-bit Integer	R																											
2	forum.title	String	R																											
3	forum.creationDate	DateTime	R																											
4	person.id	64-bit Integer	R																											
5	postCount	32-bit Integer	A																											
sort	<table><tr><td>1</td><td>postCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>forum.id</td><td>↑</td><td></td></tr></table>					1	postCount	↓		2	forum.id	↑																		
1	postCount	↓																												
2	forum.id	↑																												
limit	20																													
CPs	1.1, 1.2, 1.4, 2.1, 2.2, 2.4, 3.3																													

BI / read / 5

query	BI / read / 5																													
title	Top posters in a country																													
pattern	<div><pre>graph TD Country[Country] -- isPartOf --> City[City] City -- isLocatedIn --> member[member: Person] member -- hasMember --> forum[forum: Forum] forum -- hasMember --> person[person: Person] forum -- containerOf --> post[post: Post] person -- hasCreator --> post subgraph top100 [top 100] Country City member end subgraph count post end Forum[Forum] -- is one of the top 100 forums</pre></div>																													
desc.	<p>Find the most popular <i>Forums</i> for a given <i>Country</i>, where the popularity of a <i>Forum</i> is measured by the number of members that <i>Forum</i> has from the given <i>Country</i>.</p> <p>For each member of the 100 most popular <i>Forums</i>, count the number of <i>Posts</i> they made in any of those (most popular) <i>Forums</i>.</p>																													
params	<div><div>1</div><div>country</div><div>String</div><div></div></div>																													
result	<table><tr><td>1</td><td>person.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>person.firstName</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>person.lastName</td><td>String</td><td>R</td><td></td></tr><tr><td>4</td><td>person.creationDate</td><td>DateTime</td><td>R</td><td></td></tr><tr><td>5</td><td>postCount</td><td>32-bit Integer</td><td>A</td><td></td></tr></table>					1	person.id	64-bit Integer	R		2	person.firstName	String	R		3	person.lastName	String	R		4	person.creationDate	DateTime	R		5	postCount	32-bit Integer	A	
1	person.id	64-bit Integer	R																											
2	person.firstName	String	R																											
3	person.lastName	String	R																											
4	person.creationDate	DateTime	R																											
5	postCount	32-bit Integer	A																											
sort	<table><tr><td>1</td><td>postCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>person.id</td><td>↑</td><td></td></tr></table>					1	postCount	↓		2	person.id	↑																		
1	postCount	↓																												
2	person.id	↑																												
limit	100																													
CPs	1.2, 1.4, 1.5, 2.1, 2.2, 2.3, 2.4, 3.3, 5.3, 6.1																													

BI / read / 6

query	BI / read / 6																													
title	Most active Posters of a given Topic																													
pattern																														
desc.	<p>Get <i>Persons</i> who have created a <i>Message</i> (<i>Post</i> or <i>Comment</i>) with a given <i>Tag</i>. Each <i>Person</i> has a score, computed as follows:</p> <ul style="list-style-type: none">Count of <i>Messages</i> with the given <i>Tag</i> (messageCount).Count of <i>Likes</i> (likeCount) and <i>Comments</i> (replyCount) in reply of their <i>Messages</i> with the given <i>Tag</i> (direct relation not transitive). <p>The score is calculated as a sum with the following weights:</p> <ul style="list-style-type: none"><i>Messages</i> (messageCount) are multiplied by 1,<i>Comments</i> to <i>Messages</i> (replyCount) are multiplied by 2,<i>Likes</i> (likeCount) are multiplied by 10.																													
params	<table><tr><td>1</td><td>tag</td><td>String</td><td></td></tr></table>					1	tag	String																						
1	tag	String																												
result	<table><tr><td>1</td><td>person.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>replyCount</td><td>32-bit Integer</td><td>A</td><td></td></tr><tr><td>3</td><td>likeCount</td><td>32-bit Integer</td><td>A</td><td></td></tr><tr><td>4</td><td>messageCount</td><td>32-bit Integer</td><td>A</td><td></td></tr><tr><td>5</td><td>score</td><td>32-bit Integer</td><td>A</td><td></td></tr></table>					1	person.id	64-bit Integer	R		2	replyCount	32-bit Integer	A		3	likeCount	32-bit Integer	A		4	messageCount	32-bit Integer	A		5	score	32-bit Integer	A	
1	person.id	64-bit Integer	R																											
2	replyCount	32-bit Integer	A																											
3	likeCount	32-bit Integer	A																											
4	messageCount	32-bit Integer	A																											
5	score	32-bit Integer	A																											
sort	<table><tr><td>1</td><td>score</td><td>↓</td><td></td></tr><tr><td>2</td><td>person.id</td><td>↑</td><td></td></tr></table>					1	score	↓		2	person.id	↑																		
1	score	↓																												
2	person.id	↑																												
limit	100																													
CPs	1.2, 2.3																													

BI / read / 7

query	BI / read / 7			
title	Most authoritative users on a given topic			
pattern	<pre> graph LR Tag[Tag: name = \$tag] M1[Message] -- hasTag --> Tag M2[Message] -- hasTag --> Tag M1 -- hasCreator --> P1[Person: person: Person, id] M2 -- hasCreator --> P1 P2[Person] -- likes --> M2 P3[Person] -- hasCreator --> M3[Message] subgraph AuthorityScoreCalc [authorityScore = count] M3 -- likes --> P3 end </pre>			
desc.	<p>Given a <i>Tag</i>, find all <i>Persons</i> that ever created a <i>Message</i> with the given <i>Tag</i>. For each of these <i>Persons</i> compute their “authority score” as follows:</p> <ul style="list-style-type: none"> • The “authority score” is the sum of “popularity scores” of the <i>Persons</i> that liked any of that <i>Person</i>’s <i>Messages</i> with the given <i>Tag</i>. • A <i>Person</i>’s “popularity score” is defined as the total number of likes on all of their <i>Mes-sages</i>. 			
params	1	tag	String	
result	1	person1.id	64-bit Integer	R
	2	authorityScore	32-bit Integer	A
sort	1	authorityScore	↓	
	2	person1.id	↑	
limit	100			
CPs	1.2, 2.3, 3.2, 3.3, 6.1			

BI / read / 8

query	BI / read / 8														
title	Related topics														
pattern	<pre>graph LR Message[Message] -- hasTag --> Tag1[tag: Tag name = \$tag] Comment[comment: Comment] -- replyOf* --> Message Comment -- hasTag --> Tag2[relatedTag: Tag name != \$tag name] Tag1 -.-> hasTag Comment subgraph CountBox [count] Comment end</pre>														
desc.	Find all <i>Messages</i> that have a given <i>Tag</i> . Find the related <i>Tags</i> attached to replies of these <i>Mes-</i> <i>sages</i> (direct relation not transitive), but only of those replies that do not have the given <i>Tag</i> . Group the <i>Tags</i> by name, and get the count of replies in each group.														
params	<table><tr><td>1</td><td>tag</td><td>String</td><td></td></tr></table>					1	tag	String							
1	tag	String													
result	<table><tr><td>1</td><td>relatedTag.name</td><td>String</td><td>R</td><td></td></tr><tr><td>2</td><td>count</td><td>32-bit Integer</td><td>A</td><td></td></tr></table>					1	relatedTag.name	String	R		2	count	32-bit Integer	A	
	1	relatedTag.name	String	R											
2	count	32-bit Integer	A												
sort	<table><tr><td>1</td><td>count</td><td>↓</td><td></td></tr><tr><td>2</td><td>relatedTag.name</td><td>↑</td><td></td></tr></table>					1	count	↓		2	relatedTag.name	↑			
	1	count	↓												
2	relatedTag.name	↑													
limit	100														
CPs	1.6, 3.3, 5.2														

BI / read / 9

query	BI / read / 9			
title	Forum with related Tags			
pattern				
desc.	<p>Given two <i>TagClasses</i> (<i>tagClass1</i> and <i>tagClass2</i>), find <i>Forums</i> that contain at least one <i>Post</i> with a <i>Tag</i> from <i>tagClass1</i> and at least one <i>Post</i> with a <i>Tag</i> from <i>tagClass2</i> (direct children not transitive) – this may be the same <i>Post</i>.</p> <p>Consider the <i>Forums</i> with a number of members greater than a given threshold. For every such <i>Forum</i>, count the number of <i>Posts</i> that have a <i>Tag</i> from <i>TagClass1</i> (<i>count1</i>), and the number of <i>Posts</i> that have a <i>Tag</i> from <i>TagClass2</i> (<i>count2</i>).</p>			
params	1	tagClass1	String	
	2	tagClass2	String	
	3	threshold	32-bit Integer	
result	1	forum.id	64-bit Integer	R
	2	count1	32-bit Integer	A
	3	count2	32-bit Integer	A
sort	1	abs(count2 - count1)	↓	
	2	forum.id	↑	
limit	100			
CPs	1.2, 1.4, 2.1, 2.3, 2.4			

BI / read / 10

query	BI / read / 10			
title	Central Person for a Tag			
pattern				
desc.	<p>Given a <i>Tag</i>, find all <i>Persons</i> that are either interested in the <i>Tag</i>, or have written a <i>Message</i> (<i>Post</i> or <i>Comment</i>) with a creation date after a given date and that has a given <i>Tag</i>. For each <i>Person</i>, compute the score as the sum of the following two aspects:</p> <ul style="list-style-type: none"> • 100, if the <i>Person</i> has this <i>Tag</i> as their interest, or 0 otherwise • number of <i>Messages</i> by this <i>Person</i> with the given <i>Tag</i> <p>Also, for each <i>Person</i>, compute the sum of the score of the <i>Person</i>'s friends (<i>friendsScore</i>).</p>			
params	1	tag	String	
	2	date	Date	
result	1	person.id	64-bit Integer	R
	2	score	32-bit Integer	A
	3	friendsScore	32-bit Integer	A
	The sum of the score of the Person's friends			
sort	1	score + friendsScore	↓	
	2	person.id	↑	
limit	100			
CPs	1.2, 2.1, 2.3, 3.2			

BI / read / 11

query	BI / read / 11			
title	Unrelated replies			
pattern				
desc.	Find those <i>Persons</i> of a given <i>Country</i> that replied to any <i>Message</i> , such that the reply does not have any <i>Tag</i> in common with the <i>Message</i> (only direct replies are considered, transitive ones are not). Consider only those replies not containing any word from a given <code>blacklist</code> . For each <i>Person</i> and valid reply, retrieve the <i>Tags</i> associated with the reply, and retrieve the number of likes on the reply.			
params	1	country	String	
	2	blacklist	String[]	
result	1	person.id	64-bit Integer	R
	2	tag.name	String	R
	3	countLikes	32-bit Integer	A
	4	countReplies	32-bit Integer	A
				The count of Likes to replies with that Tag.
				The count of replies with that Tag.
sort	1	countLikes	↓	
	2	person.id	↑	
	3	tag.name	↑	
limit	100			
CPs	1.1, 2.1, 2.2, 2.3, 3.1, 3.2, 6.1			

BI / read / 12

query	BI / read / 12																													
title	Trending Posts																													
pattern	<pre>graph TD subgraph MessageBox [message: Message] direction TB F["creationDate > \$date"] ID[id] end subgraph CreatorBox [creator: Person] direction TB FN[firstName] LN[lastName] end subgraph LikesBox [likes] direction TB P[Person] end MessageBox -- hasCreator --> CreatorBox LikesBox --> MessageBox P -- count --> LikesBox</pre>																													
desc.	Find all <i>Messages</i> created after a given date, that received more than a given number of likes (likeThreshold).																													
params	<table><tr><td>1</td><td>date</td><td>Date</td><td></td></tr><tr><td>2</td><td>likeThreshold</td><td>32-bit Integer</td><td></td></tr></table>	1	date	Date		2	likeThreshold	32-bit Integer																						
1	date	Date																												
2	likeThreshold	32-bit Integer																												
result	<table><tr><td>1</td><td>message.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>message.creationDate</td><td>DateTime</td><td>R</td><td></td></tr><tr><td>3</td><td>creator.firstName</td><td>String</td><td>R</td><td>The first name of the post's creator</td></tr><tr><td>4</td><td>creator.lastName</td><td>String</td><td>R</td><td>The last name of the post's creator</td></tr><tr><td>5</td><td>likeCount</td><td>32-bit Integer</td><td>A</td><td>The number of Likes the Post received</td></tr></table>	1	message.id	64-bit Integer	R		2	message.creationDate	DateTime	R		3	creator.firstName	String	R	The first name of the post's creator	4	creator.lastName	String	R	The last name of the post's creator	5	likeCount	32-bit Integer	A	The number of Likes the Post received				
1	message.id	64-bit Integer	R																											
2	message.creationDate	DateTime	R																											
3	creator.firstName	String	R	The first name of the post's creator																										
4	creator.lastName	String	R	The last name of the post's creator																										
5	likeCount	32-bit Integer	A	The number of Likes the Post received																										
sort	<table><tr><td>1</td><td>likeCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>message.id</td><td>↑</td><td></td></tr></table>	1	likeCount	↓		2	message.id	↑																						
1	likeCount	↓																												
2	message.id	↑																												
limit	100																													
CPs	1.2, 2.2, 3.1, 6.1																													

BI / read / 13

query	BI / read / 13				
title	Popular Tags per month in a country				
pattern	<div><div>Country</div><div>name = \$country</div></div> <div><div>Message</div><div>year(creationDate) month(creationDate)</div></div> <div><div>tag: Tag</div></div> <div>isLocatedIn</div> <div>hasTag</div>				
desc.	Find all <i>Messages</i> in a given <i>Country</i> , as well as their <i>Tags</i> . Group <i>Messages</i> by creation year and month. For each group, find the 5 most popular <i>Tags</i> , where popularity is the number of <i>Messages</i> (from within the same group) where the <i>Tag</i> appears.				
params	<div><div>1</div><div>country</div><div>String</div><div></div></div>				
result	<div><div><div>1</div><div>year</div><div>32-bit Integer</div><div>C</div><div>year(message.creationDate)</div></div><div><div>2</div><div>month</div><div>32-bit Integer</div><div>C</div><div>month(message.creationDate)</div></div><div><div>3</div><div>popularTags</div><div>TagPairs</div><div>C</div><div>(tag.name - String, popularity - 32-bit Integer), sorted descending by popularity, then ascending by tag name</div></div></div>				
sort	<div><div><div>1</div><div>year</div><div>↓</div><div></div></div><div><div>2</div><div>month</div><div>↑</div><div></div></div></div>				
limit	100				
CPs	1.2, 2.2, 2.3, 3.2, 6.1				

BI / read / 14

query	BI / read / 14																													
title	Top thread initiators																													
pattern	<pre>classDiagram class Person { id firstName lastName } class Message { begin creationDate end } class Message2 { begin creationDate end } Person --> Message : hasCreator Message2 --> Message : replyOf*</pre>																													
desc.	<p>For each <i>Person</i>, count the number of <i>Posts</i> they created in the time interval [begin, end] (equivalent to the number of threads they initiated) and the number of <i>Messages</i> in each of their (transitive) reply trees. When calculating <i>Message</i> counts only consider messages created within the given time interval.</p> <p>Return each <i>Person</i>, number of <i>Posts</i> they created, and the count of all <i>Messages</i> that appeared in the reply trees (including <i>Post</i> at tree root) they created.</p>																													
params	<table><tr><td>1</td><td>begin</td><td>Date</td><td></td></tr><tr><td>2</td><td>end</td><td>Date</td><td></td></tr></table>					1	begin	Date		2	end	Date																		
1	begin	Date																												
2	end	Date																												
result	<table><tr><td>1</td><td>person.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>person.firstName</td><td>String</td><td>R</td><td></td></tr><tr><td>3</td><td>person.lastName</td><td>String</td><td>R</td><td></td></tr><tr><td>4</td><td>postCount</td><td>32-bit Integer</td><td>A</td><td>The number of <i>Posts</i> created by that <i>Person</i> (the number of threads initiated)</td></tr><tr><td>5</td><td>messageCount</td><td>32-bit Integer</td><td>A</td><td>The number of <i>Messages</i> created in all the threads this <i>Person</i> initiated</td></tr></table>					1	person.id	64-bit Integer	R		2	person.firstName	String	R		3	person.lastName	String	R		4	postCount	32-bit Integer	A	The number of <i>Posts</i> created by that <i>Person</i> (the number of threads initiated)	5	messageCount	32-bit Integer	A	The number of <i>Messages</i> created in all the threads this <i>Person</i> initiated
1	person.id	64-bit Integer	R																											
2	person.firstName	String	R																											
3	person.lastName	String	R																											
4	postCount	32-bit Integer	A	The number of <i>Posts</i> created by that <i>Person</i> (the number of threads initiated)																										
5	messageCount	32-bit Integer	A	The number of <i>Messages</i> created in all the threads this <i>Person</i> initiated																										
sort	<table><tr><td>1</td><td>messageCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>person.id</td><td>↑</td><td></td></tr></table>					1	messageCount	↓		2	person.id	↑																		
1	messageCount	↓																												
2	person.id	↑																												
limit	100																													
CPs	1.2, 2.2, 2.3, 3.2, 7.2, 7.3, 7.4																													

BI / read / 15

query	BI / read / 15			
title	Social normals			
pattern				
desc.	Given a country, find all <i>Persons</i> of the <i>Country</i> whose number of friends in the given <i>Country</i> equals the (floor of) average number of friends that <i>Persons</i> of the given <i>Country</i> have in the given <i>country</i> .			
params	1	country	String	
result	1	person.id	64-bit Integer	R
	2	count	32-bit Integer	A
sort	1	person.id	↑	
limit	100			
CPs	1.2, 2.3, 3.2, 3.3, 5.3, 6.1			

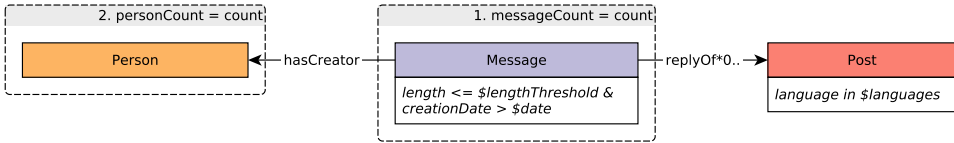
BI / read / 16

query	BI / read / 16			
title	Experts in social circle			
pattern	<pre> graph TD Country[Country] -- isPartOf --> City[City] City -- isLocatedIn --> person_Person[person: Person] person_Person -- knows* --> Person[Person] person_Person -- hasCreator --> Message[Message] Message -- hasTag --> Tag[Tag] Tag -- hasType --> TagClass[TagClass] Message -- count --> count[count] TagClass -- name = \$tagClass --> TagClass Tag -- name --> tag_Tag[tag: Tag] Person -- name = \$person --> Person City -- name = \$city --> City Country -- name = \$country --> Country </pre>			
desc.	<p>Given a <i>Person</i>, find all other <i>Persons</i> that live in a given country and are connected to given <i>Person</i> by a transitive path with length in range [minPathDistance, maxPathDistance] through the <i>knows</i> relation.</p> <p>In the path, an edge can be only traversed once while nodes can be traversed multiple times.</p> <p>For each of these <i>Persons</i>, retrieve all of their <i>Messages</i> (<i>Posts & Comments</i>) that contain at least one <i>Tag</i> belonging to a given <i>TagClass</i> (direct relation not transitive). For each <i>Message</i>, retrieve all of its <i>Tags</i>.</p> <p>Group the results by <i>Persons</i> and <i>Tags</i>, then count the <i>Messages</i> by a certain <i>Person</i> having a certain <i>Tag</i>.</p>			
params	1	personId	64-bit Integer	
	2	country	String	
	3	tagClass	String	
	4	minPathDistance	32-bit Integer	
	5	maxPathDistance	32-bit Integer	
result	1	person.id	64-bit Integer	R
	2	tag.name	String	R
	3	messageCount	32-bit Integer	A Number of <i>Messages</i> created by that <i>Person</i> containing that <i>Tag</i>
sort	1	messageCount	↓	
	2	tag.name	↑	
	3	person.id	↑	
limit	100			
CPs	1.2, 1.4, 2.3, 2.4, 3.3, 7.1, 7.2, 7.3			

BI / read / 17

query	BI / read / 17			
title	Friend triangles			
pattern				
desc.	<p>For a given country, count all the distinct triples of <i>Persons</i> such that a is friend of b, b is friend of c, and c is friend of a.</p> <p>Distinct means that given a triple t_1 in the result set R of all qualified triples, there is no triple t_2 in R such that t_1 and t_2 have the same set of elements.</p>			
params	1	country	String	
result	1	count	32-bit Integer	A
CPs	1.1, 2.3			

BI / read / 18

query	BI / read / 18															
title	How many persons have a given number of posts															
pattern																
desc.	<p>For each <i>Person</i>, count the number of <i>Messages</i> (<i>Posts</i> & <i>Comments</i>) they made. Only consider messages with:</p> <ul style="list-style-type: none">• length below the <code>lengthThreshold</code>• <code>creationDate</code> after date (exclusive, equality is not allowed)• any of the given languages <p>The language of a <i>Comment</i> is that of the <i>Post</i> that initiates the thread where <i>Comment</i> belongs to.</p>															
params	<table><tr><td>1</td><td>date</td><td>Date</td><td></td></tr><tr><td>2</td><td>lengthThreshold</td><td>32-bit Integer</td><td></td></tr><tr><td>3</td><td>languages</td><td>String[]</td><td></td></tr></table>				1	date	Date		2	lengthThreshold	32-bit Integer		3	languages	String[]	
1	date	Date														
2	lengthThreshold	32-bit Integer														
3	languages	String[]														
result	<table><tr><td>1</td><td>messageCount</td><td>32-bit Integer</td><td>A</td><td>Number of <i>Messages</i> created</td></tr><tr><td>2</td><td>personCount</td><td>32-bit Integer</td><td>A</td><td>The number of <i>Persons</i> with messageCount messages</td></tr></table>				1	messageCount	32-bit Integer	A	Number of <i>Messages</i> created	2	personCount	32-bit Integer	A	The number of <i>Persons</i> with messageCount messages		
1	messageCount	32-bit Integer	A	Number of <i>Messages</i> created												
2	personCount	32-bit Integer	A	The number of <i>Persons</i> with messageCount messages												
sort	<table><tr><td>1</td><td>personCount</td><td>↓</td><td></td></tr></table>				1	personCount	↓									
1	personCount	↓														
CPs	1.1, 1.2, 1.6, 3.2, 4.2, 4.3															

BI / read / 19

query	BI / read / 19																		
title	Stranger's interaction																		
pattern	<div><div>interactionCount = number of interactions between person and stranger, both ways</div><div><div><div><div>TagClass</div><div>name = \$tagClass1</div></div><div>hasType↑<div>Tag</div></div></div><div><div>forum1: Forum</div><div>hasTag→<div>Tag</div></div><div>hasMember→<div><div>strangersCount = count</div><div>stranger: Person</div></div></div></div><div><div>person: Person</div><div>birthday > \$date</div><div>id</div></div><div><div>forum2: Forum</div><div>hasTag→<div>Tag</div></div><div>hasMember→<div><div>strangersCount = count</div><div>stranger: Person</div></div></div></div><div><div>personA: Person</div><div>hasCreator→<div>Comment</div></div><div>replyOf→<div>Message</div></div><div>hasCreator→<div>personB: Person</div></div></div><div>interaction</div></div></div>																		
desc.	<p>For all the <i>Persons</i> born after a certain date, find all the strangers they interacted with, where strangers are <i>Persons</i> that do not <i>know</i> each other. There is no restriction on the date that strangers were born.</p> <p>Consider only strangers that are members of <i>Forums</i> tagged with a <i>Tag</i> of tagClass1 (direct children not transitive) AND members of <i>Forums</i> tagged with a <i>Tag</i> of tagClass2 (direct children not transitive). The tags may be attached to the same <i>Forum</i> or they may be attached to different <i>Forums</i>.</p> <p>Interaction is defined as follows: if <i>Person A</i> replies to a <i>Message</i> (<i>Post</i> or <i>Comment</i>) by another <i>Person B</i>, there is an “interacted with” relationship from <i>A</i> to <i>B</i>. Note that the “interacted with” relationship is directed.</p> <p>For each <i>Person</i>, count the number of strangers they interacted with and total number of times they interacted with them.</p>																		
params	<table><tr><td>1</td><td>date</td><td>Date</td><td></td></tr><tr><td>2</td><td>tagClass1</td><td>String</td><td></td></tr><tr><td>3</td><td>tagClass2</td><td>String</td><td></td></tr></table>				1	date	Date		2	tagClass1	String		3	tagClass2	String				
1	date	Date																	
2	tagClass1	String																	
3	tagClass2	String																	
result	<table><tr><td>1</td><td>person.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>strangersCount</td><td>32-bit Integer</td><td>A</td><td></td></tr><tr><td>3</td><td>interactionCount</td><td>32-bit Integer</td><td>A</td><td></td></tr></table>				1	person.id	64-bit Integer	R		2	strangersCount	32-bit Integer	A		3	interactionCount	32-bit Integer	A	
1	person.id	64-bit Integer	R																
2	strangersCount	32-bit Integer	A																
3	interactionCount	32-bit Integer	A																
sort	<table><tr><td>1</td><td>interactionCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>person.id</td><td>↑</td><td></td></tr></table>				1	interactionCount	↓		2	person.id	↑								
1	interactionCount	↓																	
2	person.id	↑																	
limit	100																		
CPs	1.1, 1.4, 2.1, 2.3, 2.4, 3.3, 5.1, 7.3, 7.4																		

BI / read / 20

query	BI / read / 20				
title	High-level topics				
pattern	<div>UNWIND \$tagClassNames[] -> \$tagClassName</div> <div><div><div>tagClass: TagClass</div><div>name = \$tagClassName</div><div>name</div></div><div>↑ isSubclassOf*0..</div><div>TagClass</div><div>↑ hasType</div><div>Tag</div><div>↑ hasTag</div><div><div>Message</div><div>count</div></div></div>				
desc.	For all given <i>TagClasses</i> , count number of <i>Messages</i> that have a <i>Tag</i> that belongs to that <i>TagClass</i> or any of its children (all descendants through a transitive relation).				
params	1	tagClasses	String[]		
result	1	tagClass.name	String	R	The <i>TagClass</i> of the root
	2	postCount	32-bit Integer	A	
sort	1	postCount	↓		
	2	tagClass.name	↑		
limit	100				
CPs	1.6, 2.1, 6.1				

BI / read / 21

query	BI / read / 21																								
title	Zombies in a country																								
pattern	<div><div><div><div>Country</div><div>name = \$country</div></div><div>City</div><div>isPartOf</div></div><div><div>zombie: Person</div><div>creationDate < \$endDate & count(message) / months < 1</div><div>isLocatedIn</div></div><div><div>Message</div><div>creationDate < \$endDate</div><div>count</div><div>hasCreator</div></div></div> <div><div><div>Person</div><div>creationDate < \$endDate</div><div>totalLikeCount = count</div><div>Person</div><div>likes</div></div><div><div>Message</div><div>zombieLikeCount = count</div><div>Person</div><div>is a zombie</div><div>likes</div></div><div><div>Person</div><div>creationDate < \$endDate</div><div>hasCreator</div></div></div>																								
desc.	<p>Find zombies within the given country, and return their zombie scores.</p> <p>A zombie is a <i>Person</i> created before the given endDate and that has created between [0, 1) Messages per month, on average, during the time range between profile creation date and the given endDate. The number of months spans the time range from creation date of the profile to the endDate and also includes partial months.</p> <p>For each <i>Person</i>, calculate the following:</p> <ul style="list-style-type: none">• zombieLikeCount: the number of likes received from zombies• totalLikeCount: the total number of likes received on that <i>Person</i>’s Messages• zombieScore: zombieLikeCount / totalLikeCount <p>In all cases, only count likes received from profiles that were created before the given endDate.</p>																								
params	<table><tr><td>1</td><td>country</td><td>String</td><td></td></tr><tr><td>2</td><td>endDate</td><td>Date</td><td></td></tr></table>					1	country	String		2	endDate	Date													
1	country	String																							
2	endDate	Date																							
result	<table><tr><td>1</td><td>person.id</td><td>64-bit Integer</td><td>R</td><td></td></tr><tr><td>2</td><td>zombieLikeCount</td><td>32-bit Integer</td><td>A</td><td></td></tr><tr><td>3</td><td>totalLikeCount</td><td>32-bit Integer</td><td>A</td><td></td></tr><tr><td>4</td><td>zombieScore</td><td>32-bit Float</td><td>A</td><td>The ratio of zombieLikeCount and totalLikeCount</td></tr></table>					1	person.id	64-bit Integer	R		2	zombieLikeCount	32-bit Integer	A		3	totalLikeCount	32-bit Integer	A		4	zombieScore	32-bit Float	A	The ratio of zombieLikeCount and totalLikeCount
1	person.id	64-bit Integer	R																						
2	zombieLikeCount	32-bit Integer	A																						
3	totalLikeCount	32-bit Integer	A																						
4	zombieScore	32-bit Float	A	The ratio of zombieLikeCount and totalLikeCount																					
sort	<table><tr><td>1</td><td>zombieScore</td><td>↓</td><td></td></tr><tr><td>2</td><td>person.id</td><td>↑</td><td></td></tr></table>					1	zombieScore	↓		2	person.id	↑													
1	zombieScore	↓																							
2	person.id	↑																							
limit	100																								
CPs	1.2, 2.1, 2.3, 2.4, 3.2, 3.3, 5.1, 5.3																								

BI / read / 22

query	BI / read / 22			
title	International dialog			
pattern	<div><div>For each pair, calculate the cost based on cases #1-5. Cases that has at least one match, add to the final score with the specified value. Each case only counts once, multiple matches do not increase to the score.</div><div><div><div>Country</div><div>name = \$country1</div></div><div>← isPartOf</div><div><div>city1: City</div><div>name</div></div><div>← isLocatedIn</div><div><div>person1: Person</div><div>id</div></div></div><div><div>Country</div><div>name = \$country2</div></div><div>← isPartOf</div><div><div>City</div></div><div>← isLocatedIn</div><div><div>person2: Person</div><div>id</div></div></div> <div><div>case 1: score += 4</div><div><div>person1: Person</div><div>likes</div><div>Message</div><div>hasCreator</div><div>person2: Person</div></div></div> <div><div>case 2: score += 1</div><div><div>person1: Person</div><div>hasCreator</div><div>Message</div><div>likes</div><div>person2: Person</div></div></div> <div><div>case 3: score += 15</div><div><div>person1: Person</div><div>knows</div><div>person2: Person</div></div></div> <div><div>case 4: score += 10</div><div><div>person1: Person</div><div>hasCreator</div><div>Comment</div><div>replyOf</div><div>Message</div><div>hasCreator</div><div>person2: Person</div></div></div> <div><div>case 5: score += 1</div><div><div>person1: Person</div><div>hasCreator</div><div>Message</div><div>replyOf</div><div>Comment</div><div>hasCreator</div><div>person2: Person</div></div></div>			
desc.	<p>Consider all pairs of people (<i>person1</i>, <i>person2</i>) such that one is located in a <i>City</i> of <i>Country</i> <i>country1</i> and the other is located in a <i>City</i> of <i>Country</i> <i>country2</i>. For each <i>City</i> of <i>Country</i> <i>country1</i>, return the highest scoring pair. The score of a pair is defined as the sum of the subscores awarded for the following kinds of interaction. The initial value is <i>score</i> = 0.</p> <div><div>1. <i>person1</i> has created a reply <i>Comment</i> to at least one <i>Message</i> by <i>person2</i>: <i>score</i> += 4</div><div>2. <i>person1</i> has created at least one <i>Message</i> that <i>person2</i> has created a reply <i>Comment</i> to: <i>score</i> += 1</div><div>3. <i>person1</i> and <i>person2</i> <i>know</i> each other: <i>score</i> += 15</div><div>4. <i>person1</i> liked at least one <i>Message</i> by <i>person2</i>: <i>score</i> += 10</div><div>5. <i>person1</i> has created at least one <i>Message</i> that was liked by <i>person2</i>: <i>score</i> += 1</div></div> <p>Consequently, the maximum score a pair can obtain is: 4 + 1 + 15 + 10 + 1 = 31</p>			
params	1	country1	String	
	2	country2	String	
result	1	person1.id	64-bit Integer	R
	2	person2.id	64-bit Integer	R
	3	city1.name	String	R
	4	score	32-bit Integer	C
sort	1	score	↓	
	2	person1.id	↑	
	3	person2.id	↑	
CPs	1.4, 1.6, 2.1, 3.1, 3.3, 5.1, 5.2, 5.3			

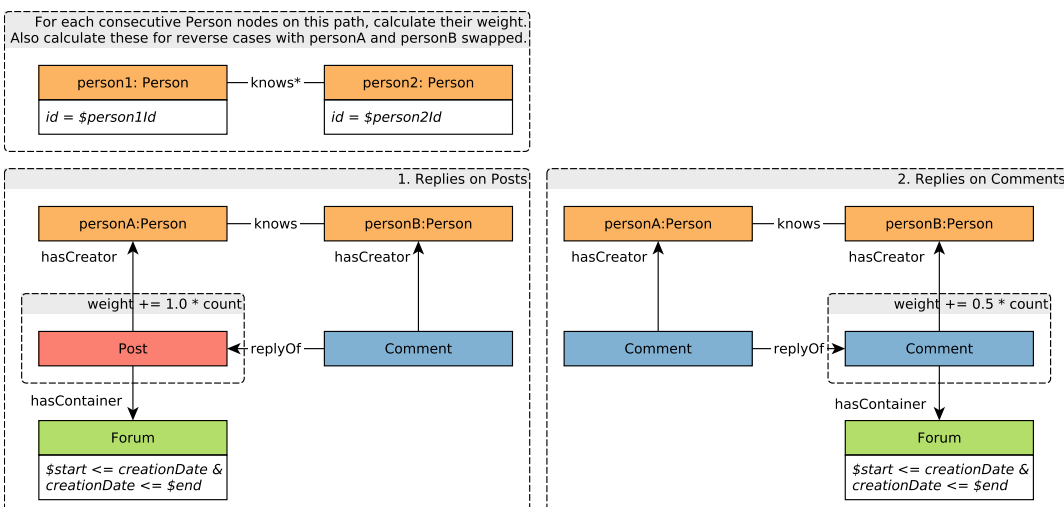
BI / read / 23

query	BI / read / 23																			
title	Holiday destinations																			
pattern																				
desc.	Count the <i>Messages</i> all residents of a given country who have written a <i>Message</i> abroad, grouped by month and Country. A <i>Message</i> was written abroad if the <i>Country</i> of the <i>Message</i> was written in is different than the <i>Country</i> of the <i>Person</i> it was written by.																			
params	<table><tr><td>1</td><td>country</td><td>String</td><td></td></tr></table>					1	country	String												
1	country	String																		
result	<table><tr><td>1</td><td>messageCount</td><td>32-bit Integer</td><td>A</td><td>The number of <i>Messages</i> in each group</td></tr><tr><td>2</td><td>country.name</td><td>String</td><td>R</td><td>The name of the destination <i>Country</i></td></tr><tr><td>3</td><td>month</td><td>32-bit Integer</td><td>C</td><td></td></tr></table>					1	messageCount	32-bit Integer	A	The number of <i>Messages</i> in each group	2	country.name	String	R	The name of the destination <i>Country</i>	3	month	32-bit Integer	C	
1	messageCount	32-bit Integer	A	The number of <i>Messages</i> in each group																
2	country.name	String	R	The name of the destination <i>Country</i>																
3	month	32-bit Integer	C																	
sort	<table><tr><td>1</td><td>messageCount</td><td>↓</td><td></td></tr><tr><td>2</td><td>country.name</td><td>↑</td><td></td></tr><tr><td>3</td><td>month</td><td>↑</td><td></td></tr></table>					1	messageCount	↓		2	country.name	↑		3	month	↑				
1	messageCount	↓																		
2	country.name	↑																		
3	month	↑																		
limit	100																			
CPs	1.6, 2.3, 2.4, 3.3, 4.3																			

BI / read / 24

query	BI / read / 24			
title	Messages by Topic and Continent			
pattern				
desc.	Find all <i>Messages</i> tagged with a <i>Tag</i> from the given <i>tagClass</i> (non-transitive). Count all <i>Messages</i> and their likes grouped by continent, year, and month.			
params	1	tagClass	String	
result	1	messageCount	32-bit Integer	A
	2	likeCount	32-bit Integer	A
	3	year	32-bit Integer	C
	4	month	32-bit Integer	C
	5	continent.name	String	R
sort	1	year	↑	
	2	month	↑	
	3	continent.name	↓	
limit	100			
CPs	1.6, 2.1, 2.3, 2.4, 3.2, 4.3			

BI / read / 25

query	BI / read / 25			
title	Weighted paths			
pattern	<p>For each consecutive Person nodes on this path, calculate their weight. Also calculate these for reverse cases with personA and personB swapped.</p> 			
desc.	<p>Given two <i>Persons</i>, find all (unweighted) shortest paths between these two <i>Persons</i>, in the sub-graph induced by the <i>knows</i> relationship.</p> <p>Then, for each path calculate a weight. The nodes in the path are <i>Persons</i>, and the weight of a path is the sum of weights between every pair of consecutive <i>Person</i> nodes in the path.</p> <p>The weight for a pair of <i>Persons</i> is calculated such that</p> <ul style="list-style-type: none"> every reply (by one of the <i>Persons</i>) to a <i>Post</i> (by the other <i>Person</i>) contributes 1.0, and every reply (by one of the <i>Persons</i>) to a <i>Comment</i> (by the other <i>Person</i>) contributes 0.5. <p>Only consider <i>Messages</i> that were created in a <i>Forum</i> that was created within the timeframe [startDate, endDate].</p> <p>Return all the paths with shortest length, and their weights.</p>			
params	1	person1Id	64-bit Integer	
	2	person2Id	64-bit Integer	
	3	startDate	Date	
	4	endDate	Date	
result	1	person.id	64-bit Integer[]	R Identifiers representing an ordered sequence of the <i>Persons</i> in the path weight
sort	1	weight	↓	The order of paths with the same weight is unspecified
CPs	1.2, 2.1, 2.2, 2.4, 3.3, 5.1, 5.3, 7.2, 7.3			

5.2 Update Query Descriptions

TODO

6 AUDITING RULES

This chapter describes the rules to audit benchmark runs, that is, what techniques are allowed and what are not, what must be provided to the auditor and guidelines for the auditors to perform the audit.

6.1 Preparation

The first step when doing an audit is to determine the versions of the following items that will be used for the benchmark:

- The benchmark specification
- The data generator
- The driver

These must be reported in the full disclosure report to guarantee that the benchmark run can be reproduced exactly in the future. Similarly, the test sponsor will inform the auditor the scale factor to test. Finally, a clean test system with enough space to store the scale factor must be provided, including the update streams and substitution parameters.

6.1.1 Collect System Details

The next step is to collect the technical and pricing details of the system under test. This includes the following items:

- Common name of the system, e.g. Dell PowerEdge xxxx.
- Type and number of CPUs, cores/threads per CPU, clock frequency and cache hierarchy characteristics (levels, size per level, etc.).
- The amount of the system's memory, type and frequency.
- The disk controller or motherboard type if disk controller is on the motherboard.
- For each distinct type of secondary storage device, the number and characteristics of the device.
- The number and type of network controllers.
- The number and type of network switches. Wiring must be disclosed.
- Date of availability of the system.

Only the network switches and interfaces that participate in the run need to be reported. If the benchmark execution is entirely contained on a single machine, no network need be reported. The price of the hardware in question must be disclosed and should reflect the single quantity list price that any buyer could expect when purchasing one system with the given specification. The price may be either an item by item price or a package price if the system is sold as a package

Besides hardware characteristics, also software details must be collected:

- The DBMS and operating system name and versions.
- Installation and configuration information of both the DBMS and operating system, which must be provided by the test sponsor.
- Price of the software license used, which can be tied to the number of concurrent users or size of data.
- Date of availability of the software.

Also, the test sponsor must provide all the source code relevant to the benchmark.

6.1.2 Setup the Benchmark Environment

Once all the information has been collected, the auditor will setup the environment to perform the benchmark run. This setup includes configuring the following items:

- Setup the LDBC Data generator in the test machine if datasets are not available from a trusted source.
- Setup the LDBC driver with the connectors provided by the test sponsor. The test sponsor must provide the configuration parameters to configure the driver (tcr, number of threads, etc.).

The `ldbc.snb.interactive.update_interleave` driver parameter must come from the `updateStream.properties` file, which is created by the data generator. That parameter should never be set manually. Also, make sure that the `-r1/--results_log` is enabled. Make sure that all operations are enabled and the frequencies are those for the selected scale factor. These can be found in Section B.1. If the driver will be executed on a separate machine, gather the characteristics of that machine in the same way as specified above.

6.1.3 Load Data

The test sponsor must provide all the necessary documentation and scripts to load the dataset into the database to test. The system under test must support the different data types needed by the benchmark for each of the attributes at their specified precision. No data can be filtered out, everything must be loaded. The test sponsor must provide a tool to perform arbitrary checks of the data or a shell to issue queries in a declarative language if the system supports it. The auditor will measure the time to load the data, which will be disclosed.

6.2 Running the Benchmark

Running the benchmark consists of three separate parts: (1) validating the query implementations, (2) warming the database and (3) performing the benchmark run. The queries are validated by means of the official validation datasets provided by LDBC consortium in their official software repositories. The auditor must load the provided dataset and run the driver in validation mode, which will test that the queries provide the official results.

The warmup can be performed either using the LDBC driver or externally, and the way it is performed must be disclosed.

A valid benchmark run must last at least 2 hours of simulation time (datagen time). Also, in order to be valid, a benchmark run needs to meet the following requirements. The `results_log.csv` file contains the `actual_start_time` and the `scheduled_start_time` of each of the issued queries. In order to have a valid run, 95% of the queries must meet the following condition:

$$\text{actual_start_time} - \text{scheduled_start_time} < 1 \text{ second}$$

If the execution of the benchmark is valid, the auditor must retrieve all the files from directory specified by `-rd/--results_dir` which includes configuration settings used, results log, results summary, which will be disclosed.

6.3 Recovery

Once an official run has been validated, the recovery capabilities of the system must be tested. The system and the driver must be configured in the same way as in during the benchmark execution. After a warmup period, an execution of the benchmark will be performed under the same terms as in the previous measured run.

At an arbitrary point close to 2 hours of simulation execution time, the machine will be disconnected. Then, the auditor will restart the database system and will check that the last committed update (in the driver log file) is actually in the database. The auditor will measure the time taken by the system to recover from the failure. Also, all the information about how durability is ensured must be disclosed. If checkpoints are used, these must be performed with a period of 10 minutes at most.

6.4 Serializability

Optionally, the test sponsor can execute update queries atomically. The auditor will verify that serializability is guaranteed.

7 RELATED WORK

7.1 Graph Database Benchmarks

Train Benchmark: model validation [12].

Graph database comparison [8].

An “interactive social network benchmark”: [3].

7.2 Scalable Graph Generators

The DATAGEN component is a fork of the generator described in [11].

7.3 LDBC Publications

LDBC publications are listed at <http://ldbouncil.org/publications>.

REFERENCES

- [1] R. Angles, P. A. Boncz, J. Larriba-Pey, I. Fundulaki, T. Neumann, O. Erling, P. Neubauer, N. Martínez-Bazan, V. Kotsev, and I. Toma. The Linked Data Benchmark Council: a graph and RDF industry benchmarking effort. *SIGMOD Record*, 43(1):27–31, 2014.
- [2] A. Averbuch and A. Prat-Pérez. Benchmark design for navigational pattern matching benchmarking. Technical report, Linked Data Benchmark Council, 2014. http://ldbouncil.org/sites/default/files/LDBC_D3.3.34.pdf.
- [3] S. Barahmand and S. Ghandeharizadeh. BG: A benchmark to evaluate interactive social networking actions. In *CIDR 2013, Sixth Biennial Conference on Innovative Data Systems Research, Asilomar, CA, USA, January 6-9, 2013, Online Proceedings*, 2013.
- [4] O. Erling, A. Averbuch, J. Larriba-Pey, H. Chafi, A. Gubichev, A. Prat-Pérez, M. Pham, and P. A. Boncz. The LDBC Social Network Benchmark: Interactive workload. In *SIGMOD*, pages 619–630, 2015.
- [5] G. Graefe. Query evaluation techniques for large databases. *ACM Comput. Surv.*, 25(2):73–170, 1993.
- [6] A. Iosup, T. Hegeman, W. L. Ngai, S. Heldens, A. Prat-Pérez, T. Manhardt, H. Chafi, M. Capota, N. Sundaram, M. J. Anderson, I. G. Tanase, Y. Xia, L. Nai, and P. A. Boncz. LDBC graphalytics: A benchmark for large-scale graph analysis on parallel and distributed platforms. *PVLDB*, 9(13):1317–1328, 2016.
- [7] A. Iosup, T. Hegeman, W. L. Ngai, S. Heldens, A. Prat-Pérez, T. Manhardt, H. Chafi, M. Capota, N. Sundaram, M. J. Anderson, I. G. Tanase, Y. Xia, L. Nai, and P. A. Boncz. LDBC Graphalytics: A benchmark for large-scale graph analysis on parallel and distributed platforms. *PVLDB*, 9(13):1317–1328, 2016.
- [8] M. Lissandrini, M. Brugnara, and Y. Velegrakis. An evaluation methodology and experimental comparison of graph databases. Technical report, University of Trento, 2017. <https://disi.unitn.it/~lissandrini/pdf/lissandrini-techreport-gdb.pdf>.
- [9] G. Moerkotte. Small materialized aggregates: A light weight index structure for data warehousing. In *PVLDB*, pages 476–487, 1998.
- [10] T. Neumann and G. Moerkotte. A framework for reasoning about share equivalence and its integration into a plan generator. In *BTW*, pages 7–26, 2009.
- [11] M. Pham, P. A. Boncz, and O. Erling. S3G2: A scalable structure-correlated social graph generator. In *Selected Topics in Performance Evaluation and Benchmarking - 4th TPC Technology Conference, TPCTC 2012, Istanbul, Turkey, August 27, 2012, Revised Selected Papers*, pages 156–172, 2012.
- [12] G. Szárnyas, B. Izsó, I. Ráth, and D. Varró. The Train Benchmark: Cross-technology performance evaluation of continuous model validation. *Softw Syst Model*, 2017.

A CHOKE POINTS

Introduction

Choke points are a superset of [2] with the exception of CP 7.1, which was removed and replaced with a new choke point.

Choke points X queries are displayed in Table A.1.

A.1 Aggregation Performance

CP-1.1: [QOPT] Interesting orders

TPC-H 1.2

This choke-point tests the ability of the query optimizer to exploit the interesting orders induced by some operators. Apart from clustered indexes providing key order, other operators also preserve or even induce tuple orderings. Sort-based operators create new orderings, typically the probe-side of a hash join conserves its order, etc.

Queries. BI 2 BI 4 BI 11 BI 17 BI 18 BI 19 Interactive 2 Interactive 9

CP-1.2: [QEXE] High Cardinality group-by performance

TPC-H 1.1

This choke-point tests the ability of the execution engine to parallelize group-by's with a large number of groups. Some queries require performing large group-by's. In such a case, if an aggregation produces a significant number of groups, intra query parallelization can be exploited as each thread may make its own partial aggregation. Then, to produce the result, these have to be re-aggregated. In order to avoid this, the tuples entering the aggregation operator may be partitioned by a hash of the grouping key and be sent to the appropriate partition. Each partition would have its own thread so that only that thread would write the aggregation, hence avoiding costly critical sections as well. A high cardinality distinct modifier in a query is a special case of this choke point. It is amenable to the same solution with intra query parallelization and partitioning as the group-by. We further note that scale-out systems have an extra incentive for partitioning since this will distribute the CPU and memory pressure over multiple machines, yielding better platform utilization and scalability.

Queries. BI 1 BI 2 BI 4 BI 5 BI 6 BI 7 BI 9 BI 10 BI 12 BI 13 BI 14 BI 15 BI 16
BI 18 BI 21 BI 25 Interactive 9

CP-1.3: [QEXE] Complex aggregate performance

This choke-point test the performance of the execution engine to perform complex aggregates. Many databases offer user defined aggregates and more complex aggregation operations than the basic count, sum, max and min, for example string concatenation aggregation operator. These types of aggregates are expected to benefit from the same optimizations as the basic built-in ones, for example partitioning.

Queries. Interactive 1

CP-1.4: [QOPT] Top-k push down

Top-k push down. This choke-point tests the ability of the query optimizer to perform optimizations based on top-k selections. Many times queries demand for returning the top-k elements. Once k results are obtained, extra restrictions in a selection can be added based on the properties of the kth element currently in the top-k, being more restrictive as the query advances, instead of sorting all elements and picking the highest k.

	1.1	1.2	1.3	1.4	1.5	1.6	2.1	2.2	2.3	2.4	3.1	3.2	3.3	4.1	4.2	4.3	5.1	5.2	5.3	6.1	7.1	7.2	7.3	7.4
BI 1		●										●		●										
BI 2	●	●		●			●		●		●	●												
BI 3										●	●	●		●		●			●	●				
BI 4	●	●		●			●	●		●			●											
BI 5		●		●	●		●	●	●	●			●						●	●				
BI 6		●							●															
BI 7		●							●			●	●							●				
BI 8						●							●					●						
BI 9		●		●			●		●	●														
BI 10		●					●		●			●												
BI 11	●						●	●	●		●	●								●				
BI 12		●						●			●									●				
BI 13		●						●	●			●								●				
BI 14		●						●	●			●										●	●	●
BI 15		●							●			●	●						●	●				
BI 16		●		●					●	●			●								●	●	●	
BI 17	●								●															
BI 18	●	●				●						●			●	●								
BI 19	●			●			●		●	●			●					●					●	●
BI 20						●	●													●				
BI 21		●					●		●	●		●	●					●		●				
BI 22				●		●	●				●		●					●	●	●				
BI 23						●			●	●			●				●							
BI 24						●	●		●	●		●				●								
BI 25		●					●	●		●			●					●		●			●	●
Interactive 1			●				●												●					
Interactive 2	●							●	●			●												
Interactive 3							●				●							●						
Interactive 4									●															
Interactive 5									●				●											
Interactive 6																		●						
Interactive 7								●	●			●						●						
Interactive 8										●		●	●						●					
Interactive 9	●	●						●	●			●	●											
Interactive 10									●				●	●	●			●	●		●	●		
Interactive 11				●					●	●			●											
Interactive 12					●								●									●	●	
Interactive 13													●									●	●	
Interactive 14													●									●	●	

Table A.1: Choke point X query

Queries. BI 2 BI 4 BI 5 BI 9 BI 16 BI 19 BI 22 Interactive 11

CP-1.5: [QEXE] Dependent group-by keys

TPC-H 1.4

This choke-point tests the ability of the query optimizer to exclude those functionally dependent group-bys. Sometimes queries require for group-by's on a set of columns and a key, where the value of the key determines the columns. In this situation, the aggregation operator should be able to exclude certain group-by attributes from the key matching.

Queries. BI 5 Interactive 12

CP-1.6: [QEXE] Low Cardinality group-by performance

TPC-H 1.3

This choke-point tests the ability to efficiently perform group by evaluation when only a very limited set of groups is available. This can require special strategies for parallelization, e.g. pre-aggregation when possible. This case also allows using special strategies for grouping like using array lookup if the domain of keys is small.

Queries. BI 8 BI 18 BI 20 BI 22 BI 23 BI 24

A.2 Join Performance

CP-2.1: [QOPT] Rich join order optimization

TPC-H 2.3

This choke-point tests the ability of the query optimizer to find optimal join orders. A graph can be traversed in different ways. In the relational model, this is equivalent as different join orders. The execution time of these orders may differ by orders of magnitude. Therefore, finding an efficient join (traversal) order is important, which in general, requires enumeration of all the possibilities. The enumeration is complicated by operators that are not freely re-orderable like semi-, anti-, and outer-joins. Because of this difficulty most join enumeration algorithms do not enumerate all possible plans, and therefore can miss the optimal join order. Therefore, these chokepoint tests the ability of the query optimizer to find optimal join (traversal) orders.

Queries. BI 2 BI 4 BI 5 BI 9 BI 10 BI 11 BI 19 BI 20 BI 21 BI 22 BI 24 BI 25
Interactive 1 Interactive 3

CP-2.2: [QOPT] Late projection

TPC-H 2.4

This choke-point tests the ability of the query optimizer to delay the projection of unneeded attributes until late in the execution. Queries where certain columns are only needed late in the query. In such a situation, it is better to omit them from initial table scans, as fetching them later by row-id with a separate scan operator, which is joined to the intermediate query result, can save temporal space, and therefore I/O. Late projection does have a trade-off involving locality, since late in the plan the tuples may be in a different order, and scattered I/O in terms of tuples/second is much more expensive than sequential I/O. Late projection specifically makes sense in queries where the late use of these columns happens at a moment where the amount of tuples involved has been considerably reduced; for example after an aggregation with only few unique group-by keys, or a top-k operator.

Queries. BI 4 BI 5 BI 11 BI 12 BI 13 BI 14 BI 25 Interactive 2 Interactive 7
Interactive 9

CP-2.3: [QOPT] Join type selection

This choke-point tests the ability of the query optimizer to select the proper join operator type, which implies accurate estimates of cardinalities. Depending on the cardinalities of both sides of a join, a hash or an index based join operator is more appropriate. This is especially important with column stores, where one usually has an index on everything. Deciding to use a hash join requires a good estimation of cardinalities on both the probe and build sides. In TPC-H, the use of hash join is almost a foregone conclusion in many cases, since an implementation will usually not even define an index on foreign key columns. There is a break even point between index and hash based plans, depending on the cardinality on the probe and build sides

Queries. BI 2 BI 5 BI 6 BI 7 BI 9 BI 10 BI 11 BI 13 BI 14 BI 15 BI 16 BI 17 BI 19
BI 21 BI 23 BI 24 Interactive 2 Interactive 4 Interactive 5 Interactive 7 Interactive 9
Interactive 10 Interactive 11

CP-2.4: [QOPT] Sparse foreign key joins

TPC-H 2.2

This choke-point tests the performance of join operators when the join is sparse. Sometimes joins involve relations where only a small percentage of rows in one of the tables is required to satisfy a join. When tables are larger, typical join methods can be sub-optimal. Partitioning the sparse table, using Hash Clustered indexes or implementing Bloom filter tests inside the join are techniques to improve the performance in such situations [5].

Queries. BI 3 BI 4 BI 5 BI 9 BI 16 BI 19 BI 21 BI 23 BI 24 BI 25 Interactive 8
Interactive 11

A.3 Data Access Locality**CP-3.1: [QOPT] Detecting correlation**

TPC-H 3.3

This choke-point tests the ability of the query optimizer to detect data correlations and exploiting them. If a schema rewards creating clustered indexes, the question then is which of the date or data columns to use as key. In fact it should not matter which column is used, as range- propagation between correlated attributes of the same table is relatively easy. One way is through the creation of multi-attribute histograms after detection of attribute correlation. With MinMax indexes, range-predicates on any column can be translated into qualifying tuple position ranges. If an attribute value is correlated with tuple position, this reduces the area to scan roughly equally to predicate selectivity.

Queries. BI 2 BI 3 BI 11 BI 12 BI 22 Interactive 3

CP-3.2: [STORAGE] Dimensional clustering

This chokepoint tests suitability of the identifiers assigned to entities by the storage system to better exploit data locality. A data model where each entity has a unique synthetic identifier, e.g. RDF or graph models, has some choice in assigning a value to this identifier. The properties of the entity being identified may affect this, e.g. type (label), other dependent properties, e.g. geographic location, date, position in a hierarchy etc, depending on the application. Such identifier choice may create locality which in turn improves efficiency of compression or index access.

Queries. BI 1 BI 2 BI 3 BI 7 BI 10 BI 11 BI 13 BI 14 BI 15 BI 18 BI 21 BI 24
Interactive 2 Interactive 8 Interactive 9

CP-3.3: [QEXE] Scattered index access patterns

This choke-point tests the performance of indexes when scattered accesses are performed. The efficiency of index lookup is very different depending on the locality of keys coming to the indexed access. Techniques like vectoring non-local index accesses by simply missing the cache in parallel on multiple lookups vectored on the same thread may have high impact. Also detecting absence of locality should turn off any locality dependent optimizations if these are costly when there is no locality. A graph neighborhood traversal is an example of an operation with random access without predictable locality.

Queries. BI 4 BI 5 BI 7 BI 8 BI 15 BI 16 BI 19 BI 21 BI 22 BI 23 BI 25 Interactive 5
Interactive 7 Interactive 8 Interactive 9 Interactive 10 Interactive 11 Interactive 12
Interactive 13 Interactive 14

A.4 Expression Calculation**CP-4.1: [QOPT] Common subexpression elimination**

TPC-H 4.2a

This choke-point tests the ability of the query optimizer to detect common sub-expressions and reuse their results. A basic technique helpful in multiple queries is common subexpression elimination (CSE). CSE should recognize also that average aggregates can be derived afterwards by dividing a SUM by the COUNT when those have been computed.

Queries. BI 1 BI 3 Interactive 10

CP-4.2: [QOPT] Complex boolean expression joins and selections

TPC-H 4.2d

This choke-point tests the ability of the query optimizer to reorder the execution of boolean expressions to improve the performance. Some boolean expressions are complex, with possibilities for alternative optimal evaluation orders. For instance, the optimizer may reorder conjunctions to test first those conditions with larger selectivity [9].

Queries. BI 18 Interactive 10

CP-4.3: [QEXE] Low overhead expressions interpretation

This choke-point tests the ability of efficiently evaluating simple expressions on a large number of values. A typical example could be simple arithmetic expressions, mathematical functions like floor and absolute or date functions like extracting a year.

Queries. BI 3 BI 18 BI 23 BI 24

CP-4.4: [QEXE] String matching performance

This chokepoint tests the ability of efficiently evaluating complex string matching expressions (e.g. via regular expressions).

A.5 Correlated Sub-queries**CP-5.1: [QOPT] Flattening sub-queries**

TPC-H 5.1

This choke-point tests the ability of the query optimizer to flatten execution plans when there are correlated sub-queries. Many queries have correlated sub-queries and their query plans can be flattened, such that the correlated

sub-query is handled using an equi-join, outer-join or anti-join. To execute queries well, systems need to flatten both sub-queries, the first into an equi-join plan, the second into an anti-join plan. Therefore, the execution layer of the database system will benefit from implementing these extended join variants. The ill effects of repetitive tuple-at-a-time sub-query execution can also be mitigated if execution systems by using vectorized, or block-wise query execution, allowing to run sub-queries with thousands of input parameters instead of one. The ability to look up many keys in an index in one API call creates the opportunity to benefit from physical locality, if lookup keys exhibit some clustering.

Queries. BI 19 BI 21 BI 22 BI 25 Interactive 3 Interactive 6 Interactive 7
Interactive 10

CP-5.2: [QEXE] Overlap between outer and sub-query

TPC-H 5.3

This choke-point tests the ability of the execution engine to reuse results when there is an overlap between the outer query and the sub-query. In some queries, the correlated sub-query and the outer query have the same joins and selections. In this case, a non-tree, rather DAG-shaped [10] query plan would allow to execute the common parts just once, providing the intermediate result stream to both the outer query and correlated sub-query, which higher up in the query plan are joined together (using normal query decorrelation rewrites). As such, the benchmark rewards systems where the optimizer can detect this and the execution engine supports an operator that can buffer intermediate results and provide them to multiple parent operators.

Queries. BI 8 BI 22 Interactive 10

CP-5.3: [QEXE] Intra-query result reuse

TPC-H 5.2

This choke-point tests the ability of the execution engine to reuse sub-query results when two sub-queries are mostly identical. Some queries have almost identical sub-queries, where some of their internal results can be reused in both sides of the execution plan, thus avoiding to repeat computations.

Queries. BI 3 BI 5 BI 15 BI 21 BI 22 BI 25 Interactive 1 Interactive 8

A.6 Parallelism and Concurrency

CP-6.1: [QEXE] Inter-query result reuse

TPC-H 6.3

This choke-point tests the ability of the query execution engine to reuse results from different queries. Sometimes with a high number of streams a significant amount of identical queries emerge in the resulting workload. The reason is that certain parameters, as generated by the workload generator, have only a limited amount of parameters bindings. This weakness opens up the possibility of using a query result cache, to eliminate the repetitive part of the workload. A further opportunity that detects even more overlap is the work on recycling, which does not only cache final query results, but also intermediate query results of a "high worth". Here, worth is a combination of partial-query result size, partial-query evaluation cost, and observed (or estimated) frequency of the partial-query in the workload.

Queries. BI 3 BI 5 BI 7 BI 11 BI 12 BI 13 BI 15 BI 20 Interactive 10

A.7 RDF and Graph Specifics

CP-7.1: [QEXE] Path pattern reuse

This choke-point tests the ability of the execution engine to reuse work across graph traversals. For example, when computing paths within a range of distances, it is often possible to incrementally compute longer paths by reusing paths of shorter distances that were already computed.

Queries. BI 16 Interactive 10

CP-7.2: [QOPT] Cardinality estimation of transitive paths

This choke-point tests the ability of the query optimizer to properly estimate the cardinality of intermediate results when executing transitive paths. A transitive path may occur in a “fact table” or a “dimension table” position. A transitive path may cover a tree or a graph, e.g. descendants in a geographical hierarchy vs. graph neighborhood or transitive closure in a many-to-many connected social network. In order to decide proper join order and type, the cardinality of the expansion of the transitive path needs to be correctly estimated. This could for example take the form of executing on a sample of the data in the cost model or of gathering special statistics, e.g. the depth and fan-out of a tree. In the case of hierarchical dimensions, e.g. geographic locations or other hierarchical classifications, detecting the cardinality of the transitive path will allow one to go to a star schema plan with scan of a fact table with a selective hash join. Such a plan will be on the other hand very bad for example if the hash table is much larger than the “fact table” being scanned.

Queries. BI 14 BI 16 BI 25 Interactive 12 Interactive 13 Interactive 14

CP-7.3: [QEXE] Execution of a transitive step

This choke-point tests the ability of the query execution engine to efficiently execute transitive steps. Graph workloads may have transitive operations, for example finding a shortest path between vertices. This involves repeated execution of a short lookup, often on many values at the same time, while usually having an end condition, e.g. the target vertex being reached or having reached the border of a search going in the opposite direction. For the best efficiency, these operations can be merged or tightly coupled to the index operations themselves. Also parallelization may be possible but may need to deal with a global state, e.g. set of visited vertices. There are many possible tradeoffs between generality and performance

Queries. BI 14 BI 16 BI 19 BI 25 Interactive 12 Interactive 13 Interactive 14

CP-7.4: [QEXE] Efficient evaluation of termination criteria for transitive queries

This tests the ability of a system to express termination criteria for transitive queries so that not the whole transitive relation has to be evaluated as well as efficient testing for termination.

Queries. BI 14 BI 19

B SCALE FACTOR STATISTICS

B.1 Scale Factor Statistics

Query Type	SF1	SF3	SF10	SF30	SF100	SF300	SF1000
Query 1	26	26	26	26	26	26	26
Query 2	37	37	37	37	37	37	37
Query 3	69	79	92	106	123	142	165
Query 4	36	36	36	36	36	36	36
Query 5	57	61	66	72	78	84	91
Query 6	129	172	236	316	434	580	796
Query 7	87	72	54	48	38	32	25
Query 8	45	27	15	9	5	3	1
Query 9	157	209	287	384	527	705	967
Query 10	30	32	35	37	40	44	47
Query 11	16	17	19	20	22	24	26
Query 12	44	44	44	44	44	44	44
Query 13	19	19	19	19	19	19	19
Query 14	49	49	49	49	49	49	49

Table B.1: Frequencies for each query and SF.